

# **SE3-Pose-Nets: Structured Deep Dynamics Models for Visuomotor Planning and Control**

Presenter: Kai-En Lin  
5/14/2020

# Outline

- Introduction
- Related work
  - SE3-Nets
- Algorithm
- Experiments
- Conclusion
- Future Work


# Introduction

- Problem statement
  - Observe a scene with a camera
  - Control the robot to reach a target

Camera



Scene



# Introduction

- Traditional approach
  - Data-associate the observed scene to target
    - E.g. tracking different parts of the robot
  - Model the effect of applied actions to changes to the scene
    - E.g. knowing what happens after the action
- Deep Learning approach
  - Tries to learn similar models
  - Lacks the ability to associate objects/parts across scenes

# Introduction

- Goal:
  - Devise a learning-based algorithm that allows:
    - Data-association
    - Modeling of the object dynamics
    - Correct prediction and control from the model

# Related Work

- SE3-Nets
  - Segment object parts
  - Predict  $SE(3)$  transformation for each part to target
  - No explicit modeling of data association

# Algorithm

- Given
  - an observation  $x_t$  of the scene (depth map / point cloud)
  - applied actions  $u_t$
- Predict
  - the transformed output point cloud  $x_{t+1}$

# Algorithm

- We can decompose the problem of modeling scene dynamics into:
  1. Modeling scene structure
  2. Modeling the dynamics of individual parts
  3. Combining local pose changes to model the dynamics of the entire scene
- With deep learning:
  1. An encoder to distinguish individual parts and predict a 6D pose for each of them
  2. A pose transition network to model the dynamics in the pose space. Takes source pose and action to predict the change in poses
  3. A transform layer to apply  $SE(3)$  transforms to input point cloud using predicted pose deltas



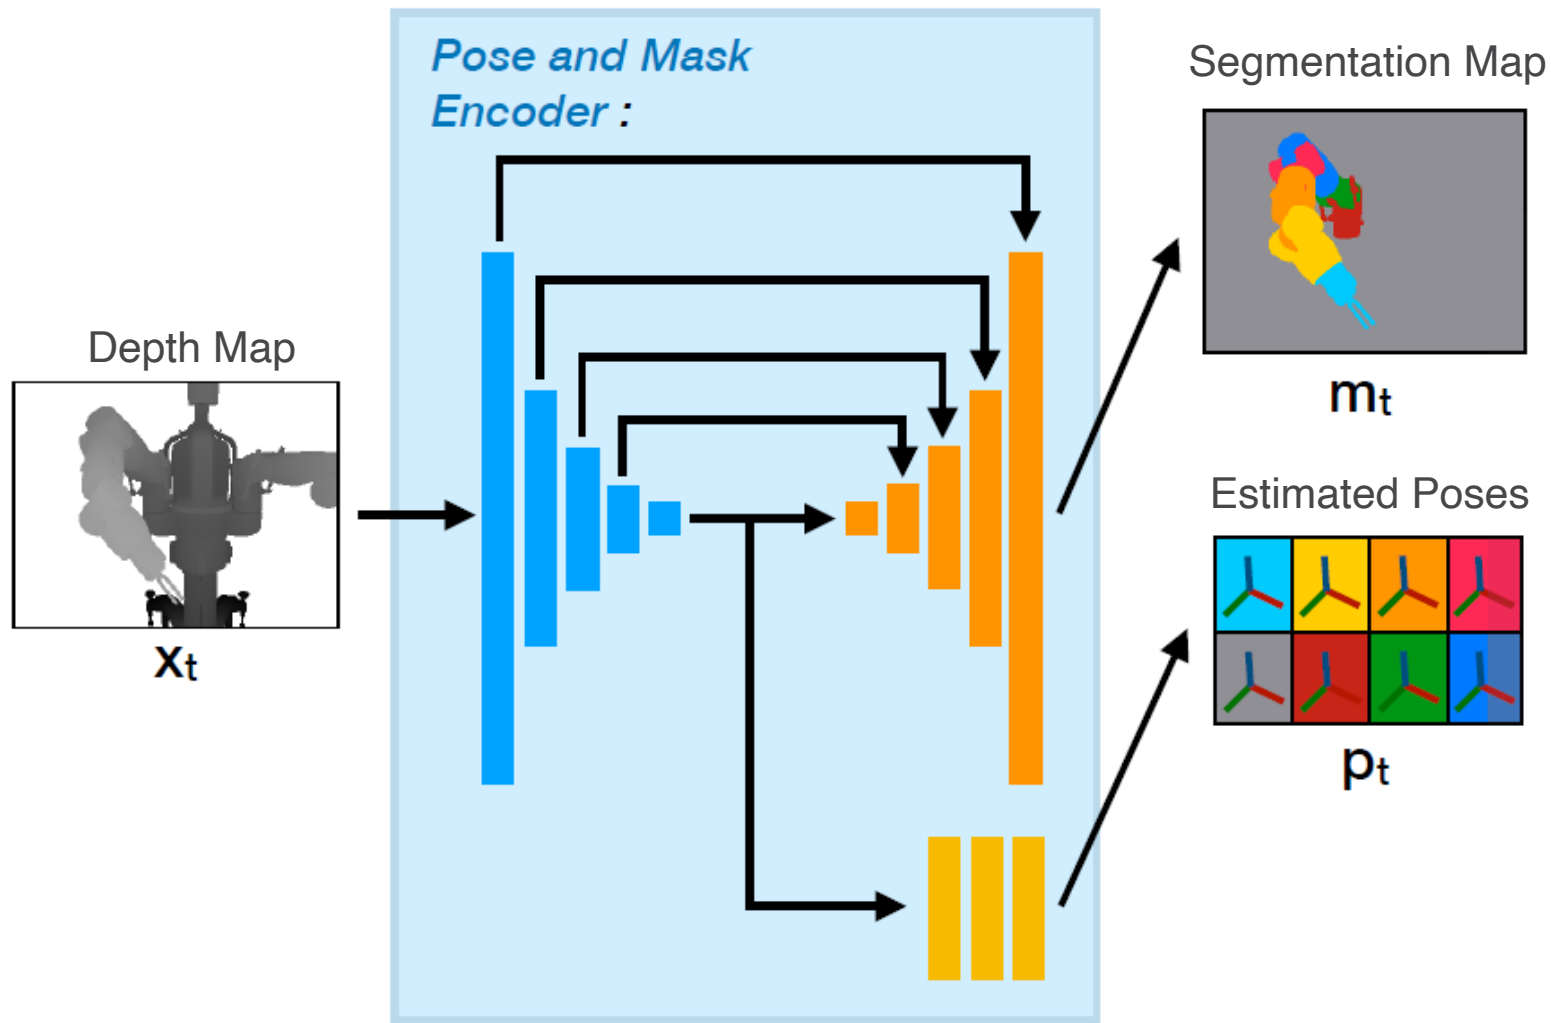
# Algorithm

- We can decompose the problem of modeling scene dynamics into:
  1. Modeling scene structure
  2. Modeling the dynamics of individual parts
  3. Combining local pose changes to model the dynamics of the entire scene
- With deep learning:
  1. An encoder to distinguish individual parts and predict a 6D pose for each of them
  2. A pose transition network to model the dynamics in the pose space. Takes source pose and action to predict the change in poses
  3. A transform layer to apply  $SE(3)$  transforms to input point cloud using predicted pose deltas

# Modeling Scene Structure

- An encoder takes the input 3D point cloud  $x_t$  and generates the following:
  - Masks for the moving parts ( $m_t$ )
  - 6D pose per segmented part ( $p_t$ )
    - 3D position
    - Orientation as 3-parameter axis-angle vector

# Modeling Scene Structure




# Algorithm

- We can decompose the problem of modeling scene dynamics into:
  1. Modeling scene structure
  2. Modeling the dynamics of individual parts
  3. Combining local pose changes to model the dynamics of the entire scene
- With deep learning:
  1. An encoder to distinguish individual parts and predict a 6D pose for each of them
  2. A pose transition network to model the dynamics in the pose space. Takes source pose and action to predict the change in poses
  3. A transform layer to apply  $SE(3)$  transforms to input point cloud using predicted pose deltas

# Modeling Part Dynamics

- A fully-connected pose transition network takes the predicted poses from the encoder ( $p_t$ ) and applied actions ( $u_t$ ) as input and predicts:
  - The change in pose ( $\Delta p_t$ ) for all  $K$  segmented parts (6D vector)

# Modeling Part Dynamics



# Algorithm

- We can decompose the problem of modeling scene dynamics into:
  1. Modeling scene structure
  2. Modeling the dynamics of individual parts
  3. Combining local pose changes to model the dynamics of the entire scene
- With deep learning:
  1. An encoder to distinguish individual parts and predict a 6D pose for each of them
  2. A pose transition network to model the dynamics in the pose space. Takes source pose and action to predict the change in poses
  3. A transform layer to apply  $SE(3)$  transforms to input point cloud using predicted pose deltas

# Predicting Scene Dynamics

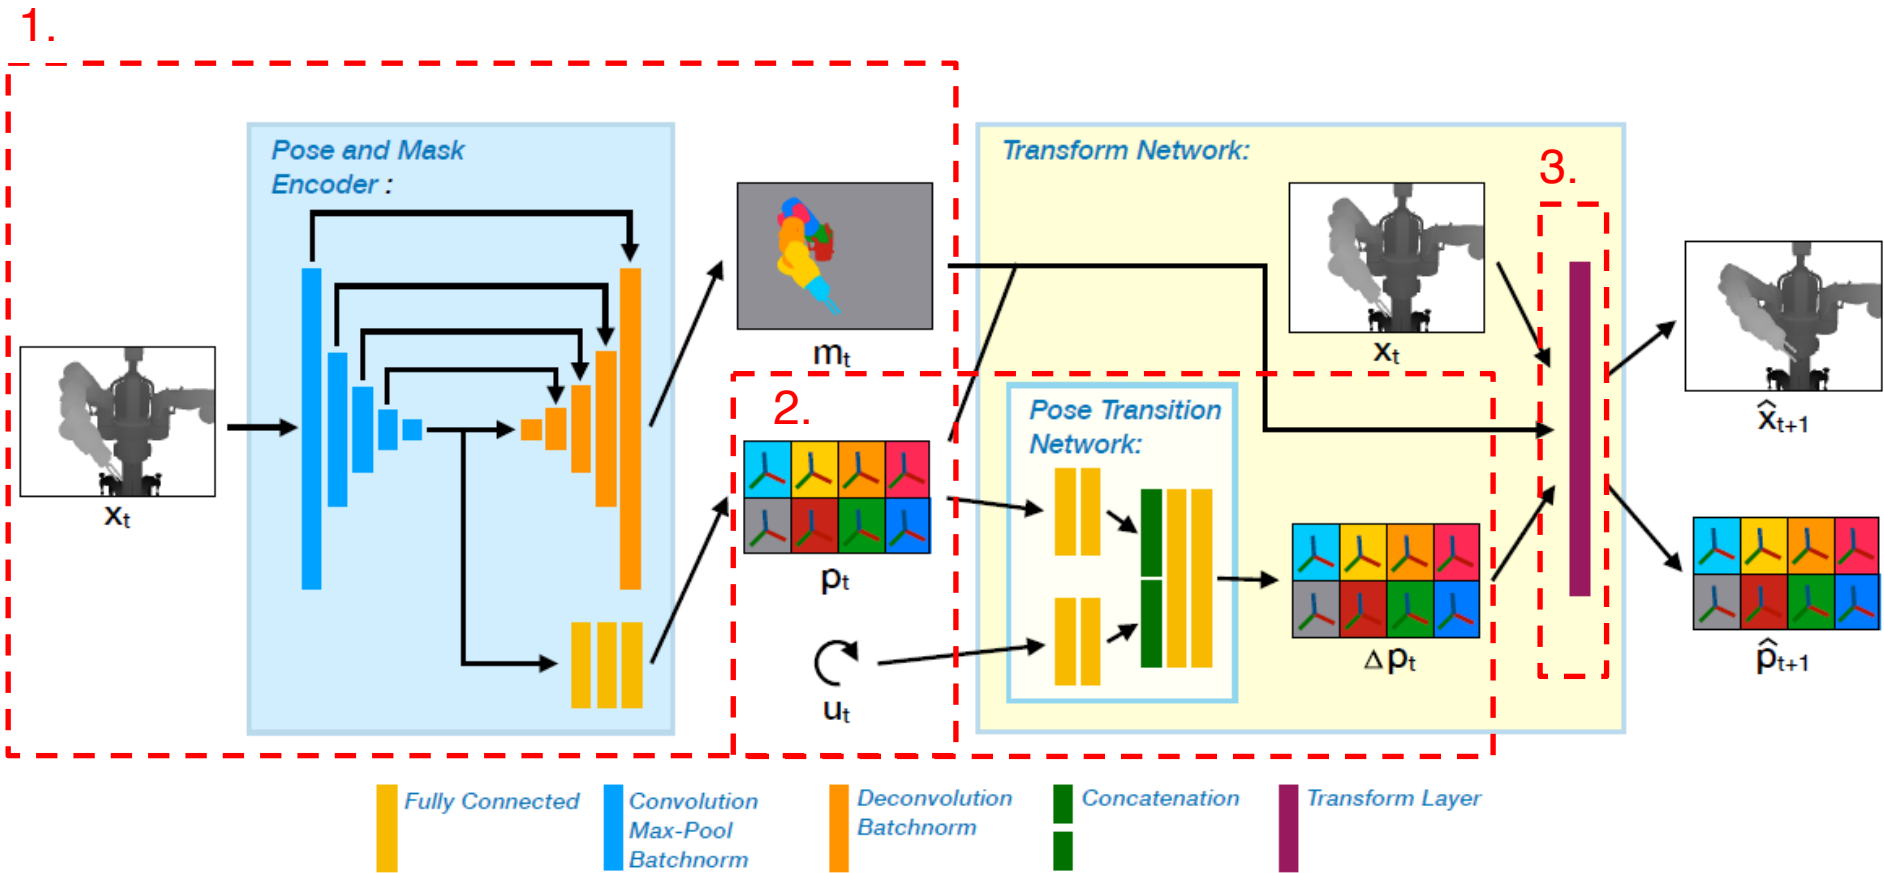
- No trainable parameters
- Given point cloud ( $x_t$ ), the predicted scene segmentation ( $m_t$ ) and the change in poses ( $\Delta p_t$ ), calculates the point cloud in the next frame ( $x_{t+1}$ ):

$$\hat{x}_{t+1}^j = \sum_{k=1}^K m_t^{kj} (R_t^k x_t^j + T_t^k),$$

where  $R_t^k$  is rotation,  $T_t^k$  translation



# Algorithm



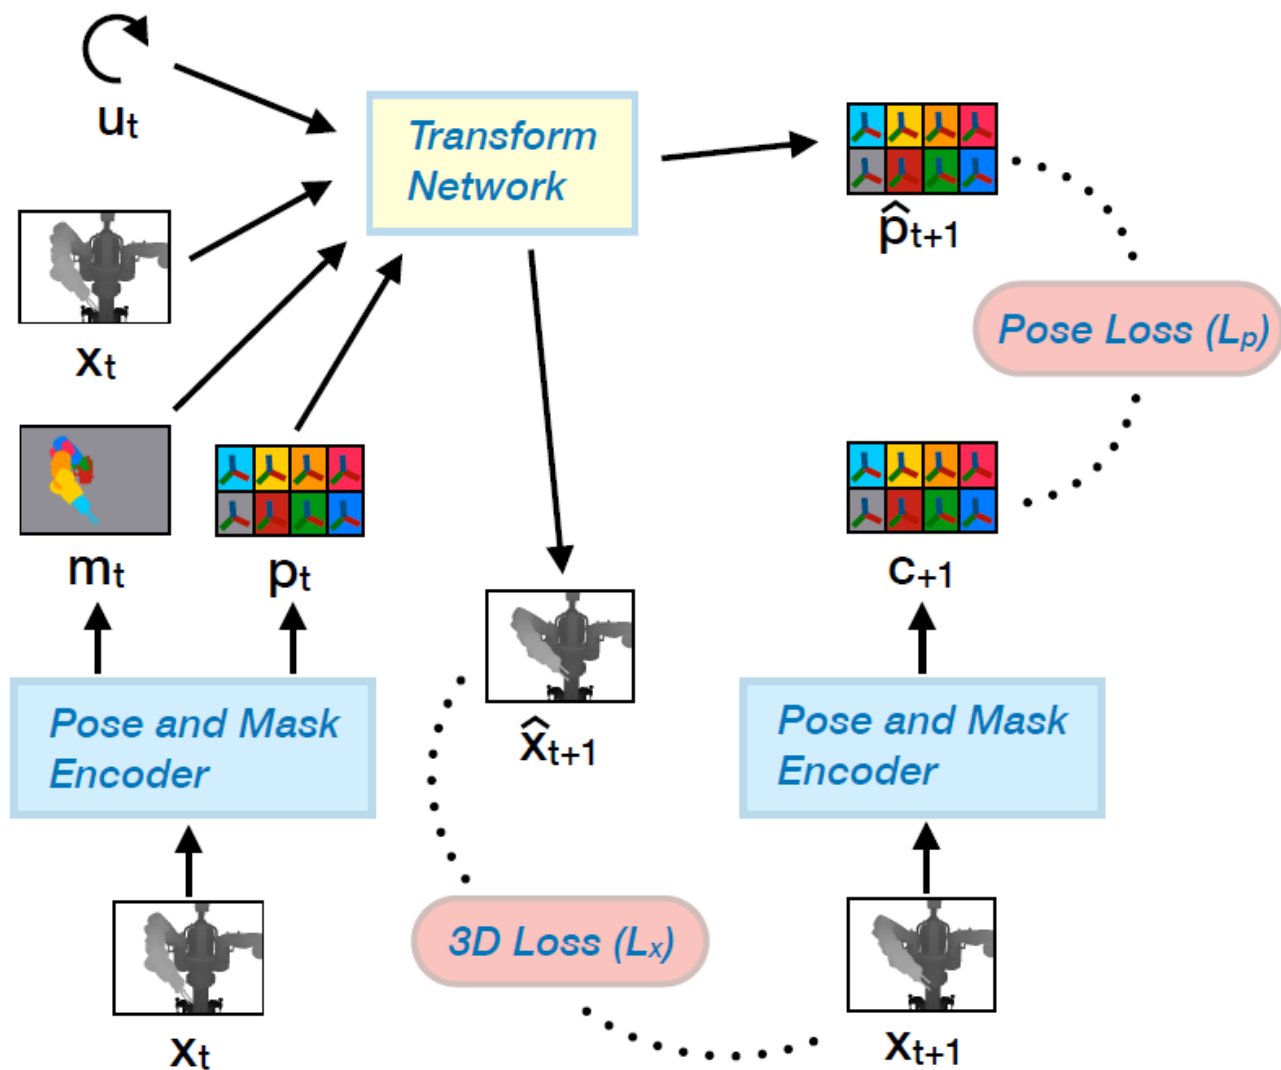
# Training

- Supervision
  - Point-wise data associations across a pair of point clouds ( $\mathbf{x}_t, \mathbf{x}_{t+1}$ )
  - Related by an action ( $\mathbf{u}_t$ )

# Training

- Total loss  $L = L_x + \gamma L_p$ 
  - 3D Loss  $L_x$
  - Pose consistency loss  $L_p$
  - $\gamma = 10$

# Training



# Training

- 3D Loss  $L_x$

$$L_x = \frac{1}{N} \sum_{i=1}^{HW} \frac{(\hat{x}_{t+1}^i - \tilde{x}_{t+1}^i)^2}{\alpha \tilde{f}^i + \beta},$$

where  $HW$  is the number of points,

$\alpha = 0.5$ ,  $\beta = 1e - 3$ ,

$\tilde{f}^i = \tilde{x}_{t+1}^i - x_t^i$ , scaling factor to make the loss scale-invariant

# Training

- Pose consistency loss  $L_p$

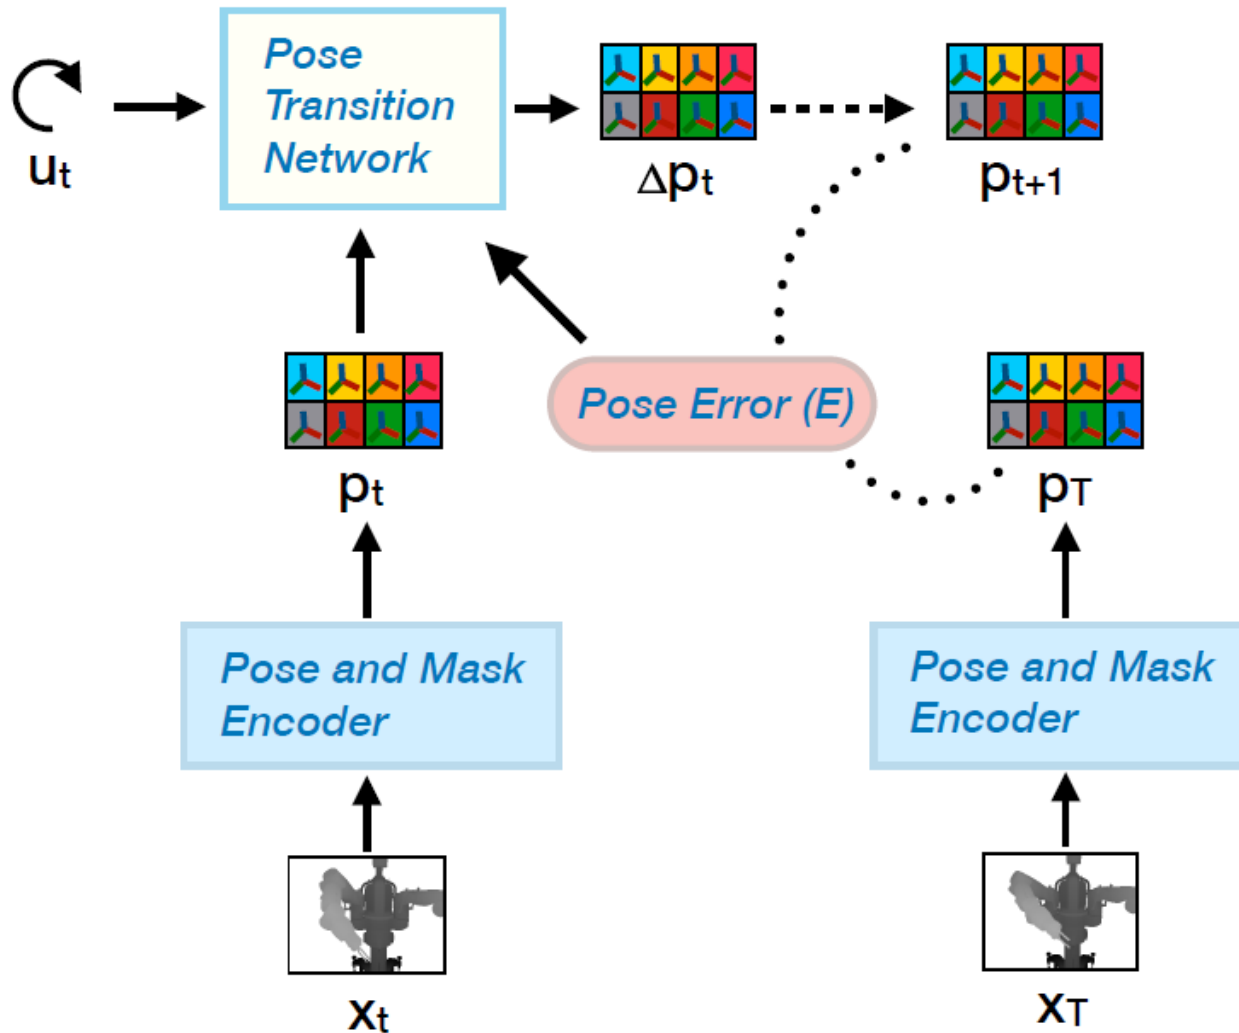
$$L_p = \frac{1}{I} \sum_{i=1}^I (\hat{p}_{t+1}^i - p_{t+1}^i)^2,$$

where  $\hat{p}_{t+1} = p_t \oplus \Delta p_t$ , the expected pose at time  $t+1$

# Closed-Loop Visuomotor Control Using SE3-Pose-Nets

- Visual servoing
  - Given the current image and the target image, generate controls to reach the target
- SE3-Pose-Nets solve this by using the latent pose space to data-associate the observations and minimizing the error between initial pose  $p_0$  and the final pose  $p_T$ .

# Closed-Loop Visuomotor Control Using SE3-Pose-Nets





# Closed-Loop Visuomotor Control Using SE3-Pose-Nets

---

**Algorithm 1** Reactive visuomotor control

---

Given: Target point cloud ( $\mathbf{x}_T$ )

Given: Pre-trained encoder ( $h_{enc}$ ) and transition model ( $h_{trans}$ )

Given: Maximum control magnitude:  $u_{max}$

Compute target pose:  $\mathbf{p}_T = h_{enc}(\mathbf{x}_T)$

**while** not converged **do**

    Receive current observation ( $\mathbf{x}_t$ )

    Predict current pose:  $\mathbf{p}_t = h_{enc}(\mathbf{x}_t)$

    Initialize control to all zeros:  $\mathbf{u}_t = 0$

    Predict change in pose:  $\Delta\mathbf{p}_t = h_{trans}(\mathbf{p}_t, \mathbf{u}_t)$

    Predict next pose:  $\hat{\mathbf{p}}_{t+1} = \mathbf{p}_t \oplus \Delta\mathbf{p}_t$

    Compute pose error:  $E = \frac{1}{I} \sum_{i=1}^I (\hat{p}_{t+1}^i - p_T^i)^2$

    Compute gradient of error w.r.t. control:  $g = \frac{dE}{d\mathbf{u}_t}$

    Compute control:  $\mathbf{u}_t = -u_{max} * \frac{g}{\|g\|}$

    Execute control  $\mathbf{u}_t$  on the robot

---

# Experiments

- SE3-Pose-Nets performs worse than other models when predicting scene dynamics
  - The pose space might make the training problem harder
  - Constraint on pose consistency is different from the prediction problem

Setting	SE3-POSE-NETS	SE3-POSE-NETS + Joint Angles	SE3-NETS	SE3-NETS + Joint Angles	Flow	Flow + Joint Angles
Simulated	0.044	0.038	0.030	<b>0.024</b>	0.035	0.030
Real	0.234	0.224	0.221	<b>0.212</b>	0.228	0.218

TABLE I: Average per-point flow MSE (cm) across tasks and networks, normalized by the number of points  $M$  that move in the ground truth data (motion magnitude  $> 1\text{mm}$ ). Our network achieves results slightly worse than the baseline networks on both simulated and real data. However, it is also solving additional tasks necessary for control.

# Experiments

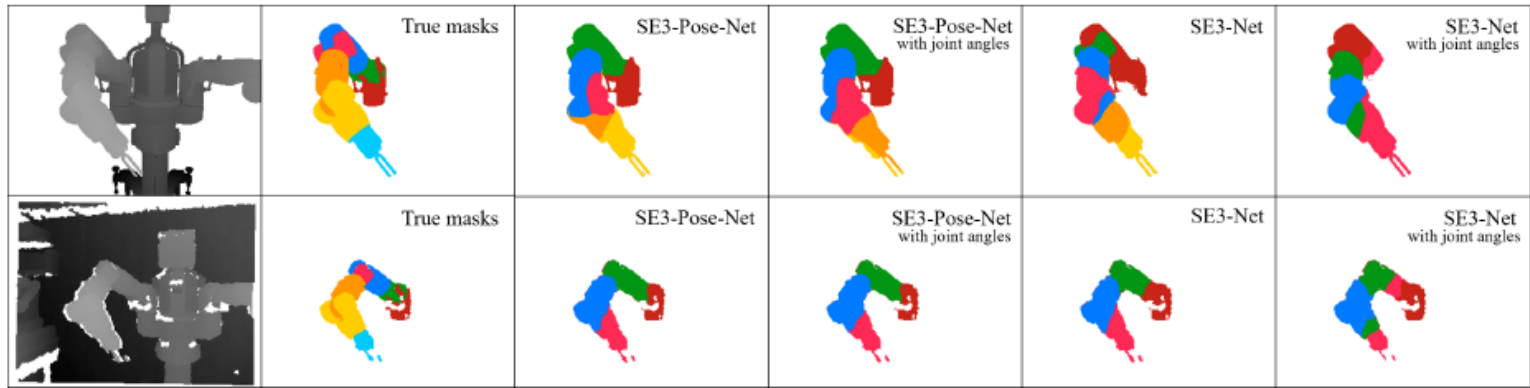


Fig. 3: Masks generated by different networks on simulated (top) and real data (bottom). From left to right: Ground truth depth, ground truth masks, masks predicted by the SE3-POSE-NET, SE3-POSE-NET with joint angles, SE3-NET and SE3-NET with joint angles.

# Experiments




Fig. 4: Convergence of joint angle error in simulated Baxter control tasks. (left): without joint angles, (middle) without joint angles and detected failure case removed (for all methods), (right) with joint angles. SE3-POSE-NETS perform as well or better than baseline methods even though baseline models have additional information in the form of ground truth-associations.

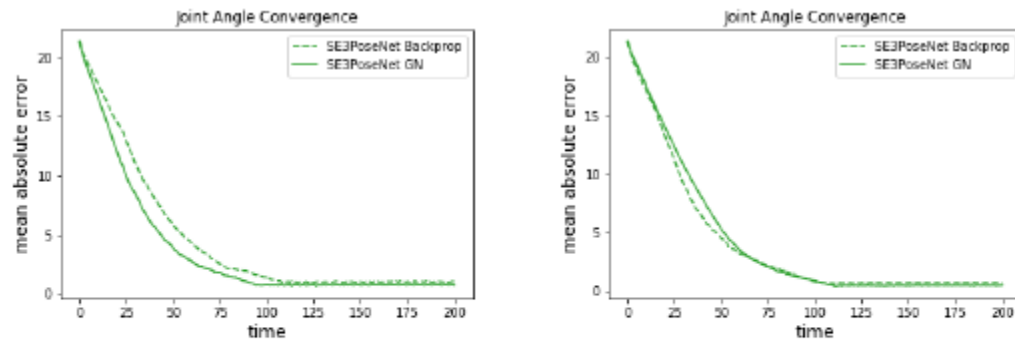
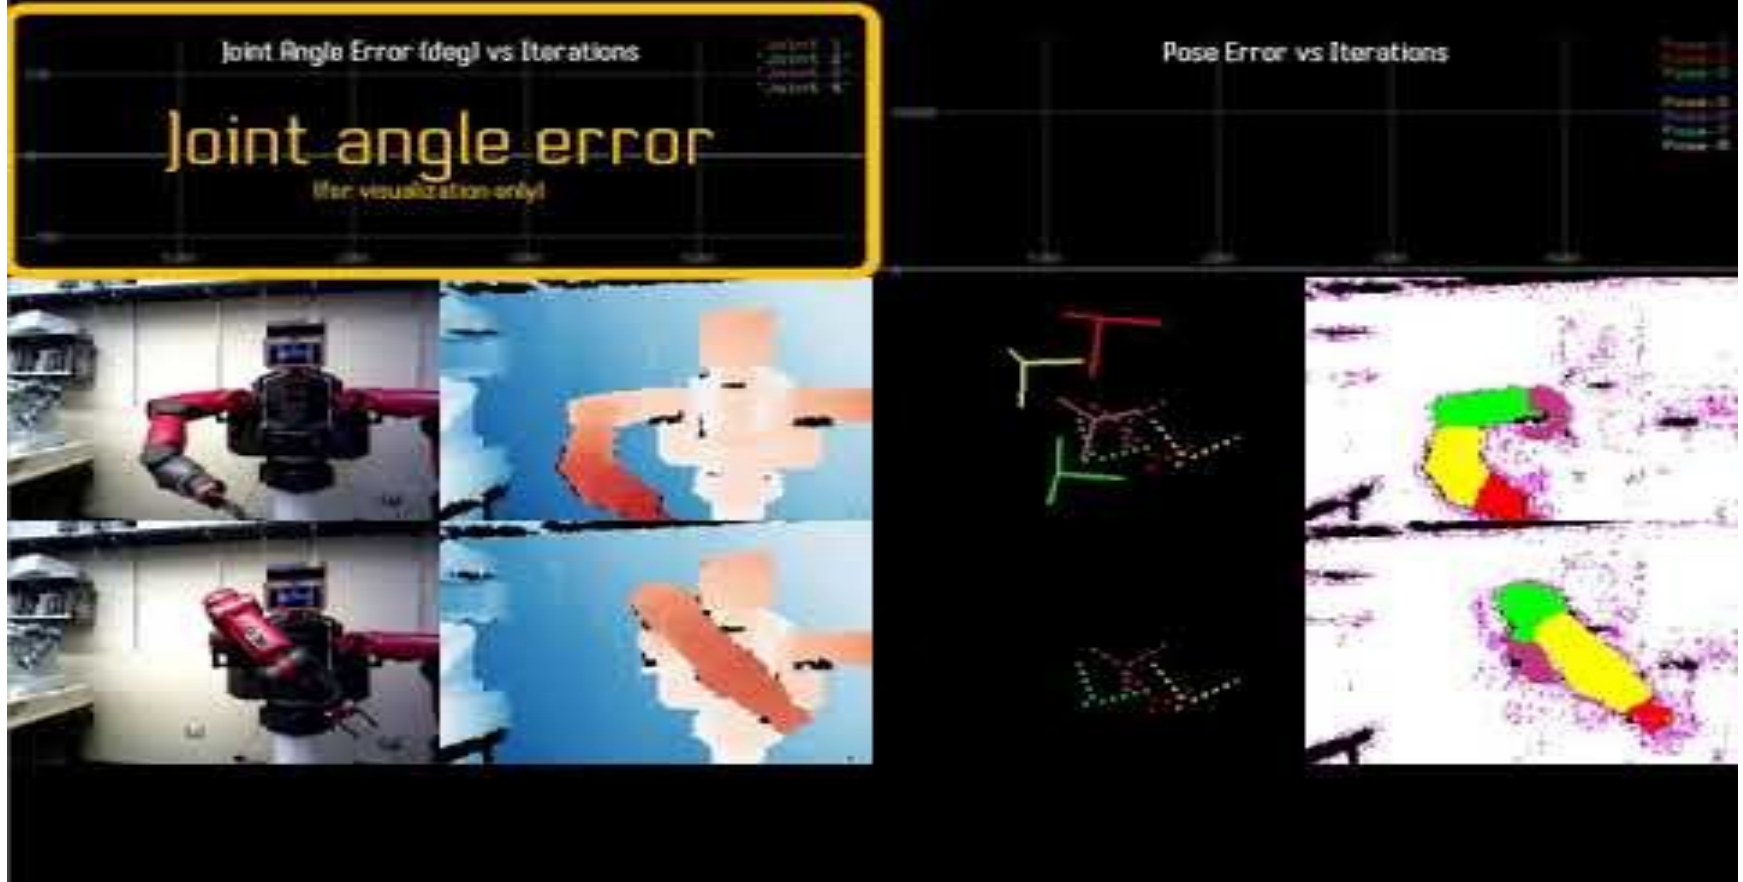


Fig. 5: Convergence of joint angle error on real Baxter control tasks (left) without joint angles (right) with joint angles (averaged across joint 0,1,2,3).

# Experiments

Example #1



# Conclusion

- SE3-Pose-Nets is an end-to-end framework for learning predictive models that enable control of objects in a scene
- It learns a consistent pose space for each individual part
- Does not require external data association
- The network enables computation of controls in the low dimensional pose space

# Future Work

- SE3-Pose-Nets has difficulties handling joints further down the kinematic chain
- Extending the system to interact with and manipulate external objects
- Long-term planning to utilize the latent pose space