

Learning Agile Robotic Locomotion Skills by Imitating Animals

Presenter: Venkatesh Prasad Venkataramanan
28th of May, 2020

Outline

- Introduction
- Related work
- Method
- Experiments
- Conclusion

Introduction

- Reproducing the diverse and agile locomotion skills of animals has been a longstanding challenge in robotics.
- Manually designed controllers can emulate many behaviours, but require effort and expertise.
- In this work, the authors present an imitation learning based approach to automate the tedious task of design of controllers



Related Work

- Performing RL training in real-world is expensive.
- Motion Imitation is successful in sim, but fails in real-world
- Focus is on domain transfer approaches i.e., sim to real world
- Here, domain adaptation is done through a method broadly classified as latent space methods.


Related Work

- Previous latent space methods used a manually designed reward function, eg ANYmal Robot
- Motion Imitation can help prevent need for manually designing reward functions.
- Motion Imitation + Latent Space Method = Success !!!


Dog Trot



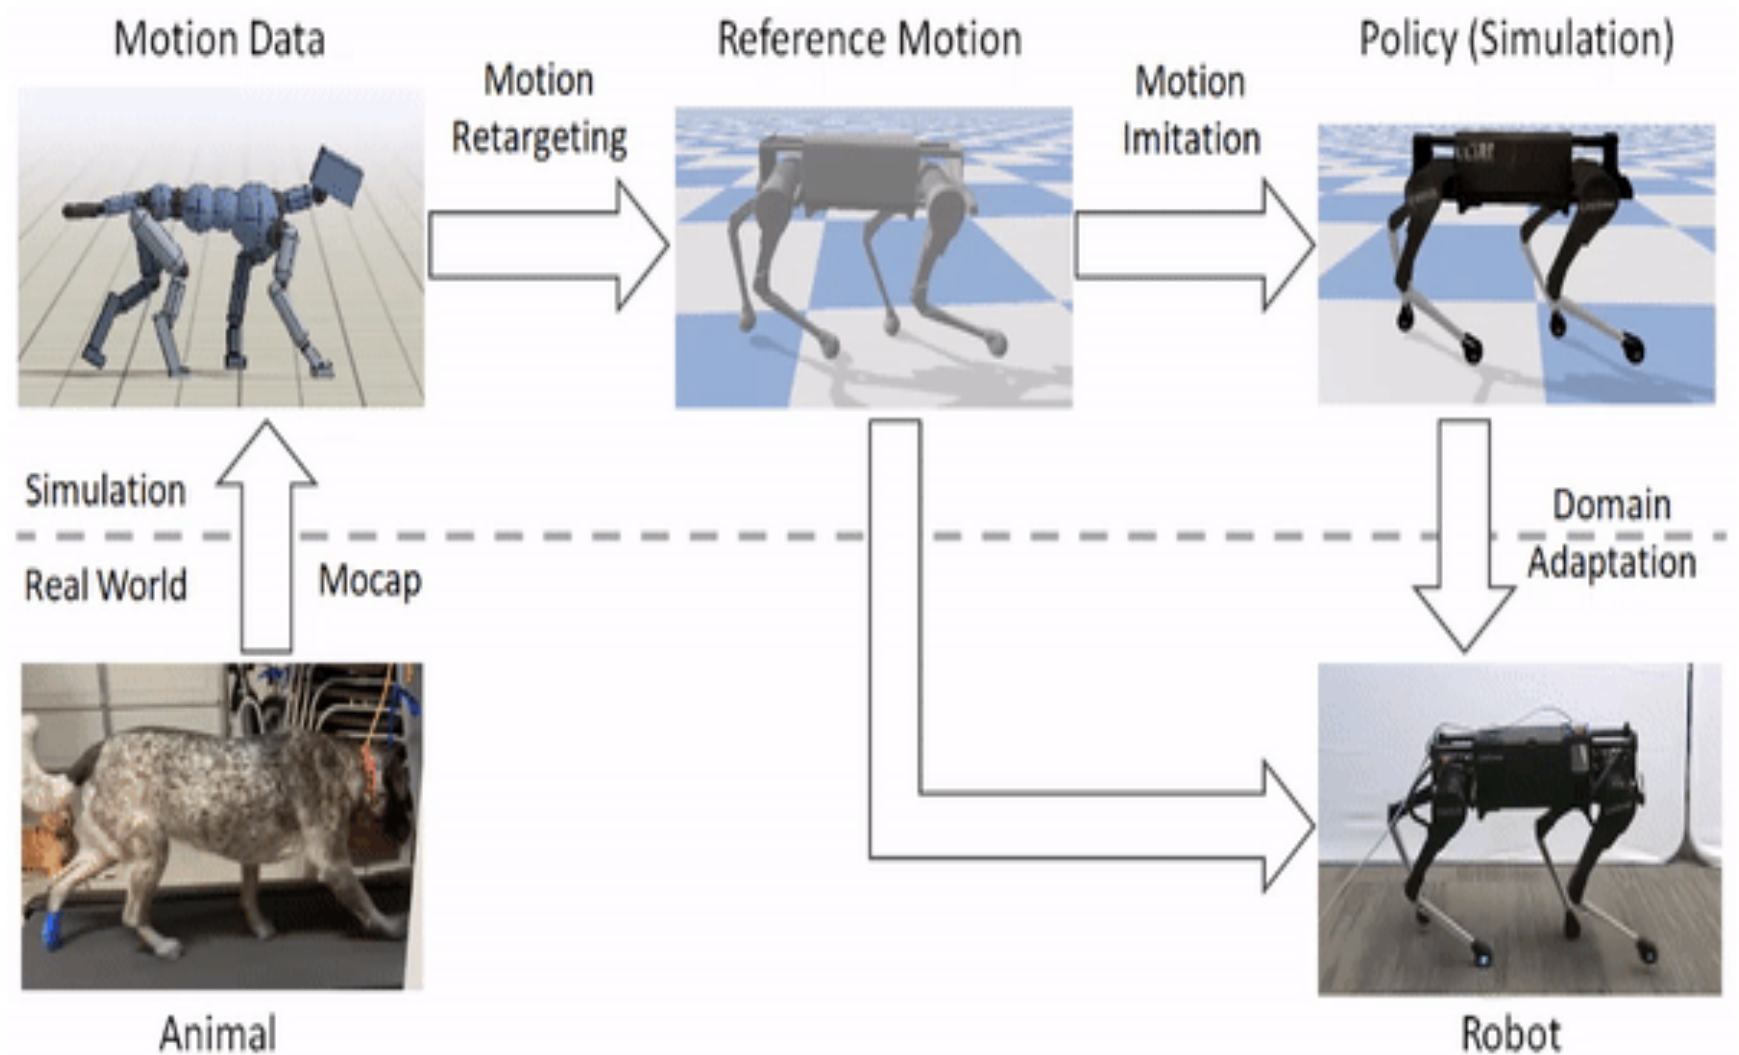
Dog Spin





Side-Steps



Overview of Method



Motion Retargeting




Motion Retargeting

- At each timestep, the source motion specifies the 3D location $\mathbf{x}_i(t)$ of each keypoint i .
- The corresponding target keypoint $\mathbf{x}_i(\mathbf{q}_t)$ is determined by the robot's pose \mathbf{q}_t
- IK is then applied to construct a sequence of poses that track the keypoints represented by $\mathbf{q}_{0:T}$

$$\arg \min_{\mathbf{q}_{0:T}} \sum_t \sum_i \|\hat{\mathbf{x}}_i(t) - \mathbf{x}_i(\mathbf{q}_t)\|^2 + (\bar{\mathbf{q}} - \mathbf{q}_t)^T \mathbf{W} (\bar{\mathbf{q}} - \mathbf{q}_t).$$

Motion Imitation



$$J(\pi) = \mathbb{E}_{\tau \sim p(\tau|\pi)} \left[\sum_{t=0}^{T-1} \gamma^t r_t \right]$$

$$p(\tau|\pi) = p(s_0) \prod_{t=0}^{T-1} p(s_{t+1}|s_t, a_t) \pi(a_t|s_t)$$

Motion Imitation

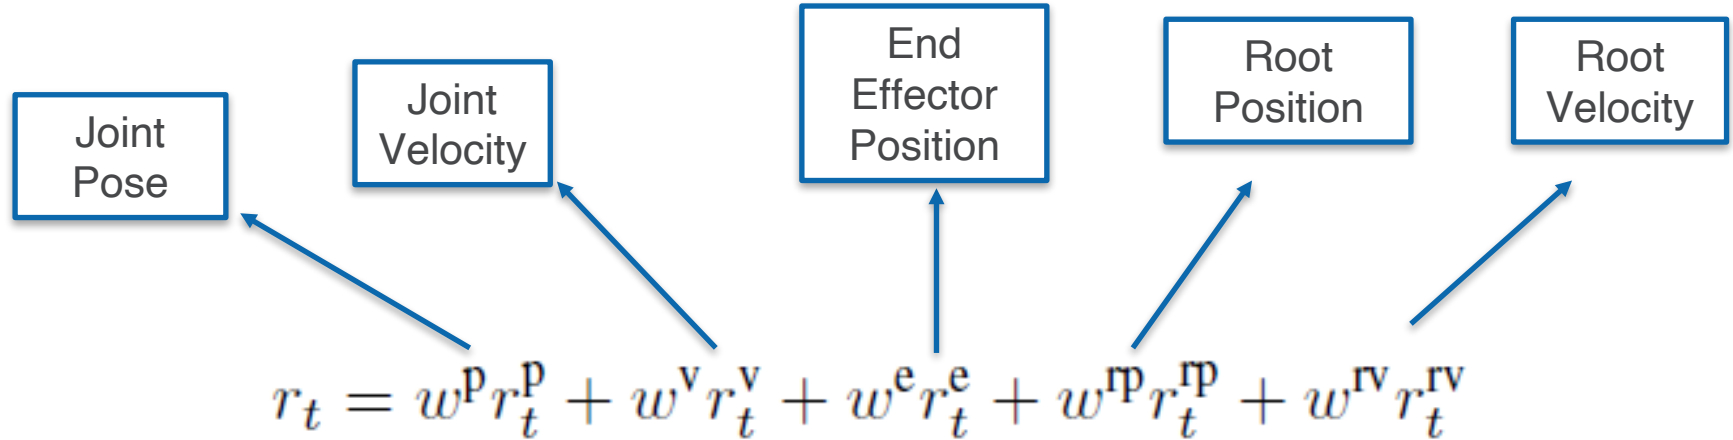
- Method used is similar to Peng et al.
- The inputs to the policy is augmented with an additional goal g_t , which specifies the motion that the robot should imitate $\pi(\mathbf{a}_t | \mathbf{s}_t, \mathbf{g}_t)$

$$\mathbf{s}_t = (\mathbf{q}_{t-2:t}, \mathbf{a}_{t-3:t-1})$$

- Pose \mathbf{q}_t taken from IMU (yaw pitch roll) and joint rotations.

Reward Function

- Again, borrowed from Peng et al.




$$w^p = 0.5, w^v = 0.05, w^e = 0.2, w^{rp} = 0.15, w^{rv} = 0.1$$

$$r_t^p = \exp \left[-5 \sum_j \|\hat{\mathbf{q}}_t^j - \mathbf{q}_t^j\|^2 \right]$$


Domain Randomization fails!!!

- Instead of training a policy in a single environment with fixed dynamics, domain randomization varies the dynamics during training
- Due to unmodeled effects in the real world, systems may nonetheless fail when deployed in a physical system.



Domain Adaptation

- Search is performed to find a latent encoding $z \sim E(z|\mu)$



Domain Adaptation

- They incorporate an information bottleneck into the encoder between the dynamics parameters \mathbf{M} and the encoding \mathbf{Z}

$$I(\mathbf{M}, \mathbf{Z}) \leq I_c.$$

$$\arg \max_{\pi, E} \mathbb{E}_{\mu \sim p(\mu)} \mathbb{E}_{\mathbf{z} \sim E(\mathbf{z}|\mu)} \mathbb{E}_{\tau \sim p(\tau|\pi, \mu, \mathbf{z})} \left[\sum_{t=0}^{T-1} \gamma^t r_t \right]$$

Domain Adaptation

$$I(\mathbf{M}, \mathbf{Z}) \leq \mathbb{E}_{\mu \sim p(\mu)} [\text{D}_{\text{KL}} [E(\cdot|\mu) || \rho(\cdot)]]$$

Stochastic
Encoder

Variational
Prior

$$\arg \max_{\pi, E} \mathbb{E}_{\mu \sim p(\mu)} \mathbb{E}_{\mathbf{z} \sim E(\mathbf{z}|\mu)} \mathbb{E}_{\tau \sim p(\tau|\pi, \mu, \mathbf{z})} \left[\sum_{t=0}^{T-1} \gamma^t r_t \right] \\ - \beta \mathbb{E}_{\mu \sim p(\mu)} [\text{D}_{\text{KL}} [E(\cdot|\mu) || \rho(\cdot)]] ,$$

Lagrange
Multiplier

Algorithm


Algorithm 1 Adaptation with Advantage-Weighted Regression

- 1: $\pi \leftarrow$ trained policy
 - 2: $\omega_0 \leftarrow \mathcal{N}(0, I)$
 - 3: $\mathcal{D} \leftarrow \emptyset$
 - 4: **for** iteration $k = 0, \dots, k_{\max} - 1$ **do**
 - 5: $\mathbf{z}_k \leftarrow$ sampled encoding from $\omega_k(\mathbf{z})$
 - 6: Rollout an episode with π conditioned \mathbf{z}_k and record the return \mathcal{R}_k
 - 7: Store $(\mathbf{z}_k, \mathcal{R}_k)$ in \mathcal{D}
 - 8: $\bar{v} \leftarrow \frac{1}{k} \sum_{i=1}^k \mathcal{R}_i$
 - 9: $\omega_{k+1} \leftarrow \arg \max_{\omega} \sum_{i=1}^k \left[\log \omega(\mathbf{z}_i) \exp \left(\frac{1}{\alpha} (\mathcal{R}_i - \bar{v}) \right) \right]$
 - 10: **end for**
-

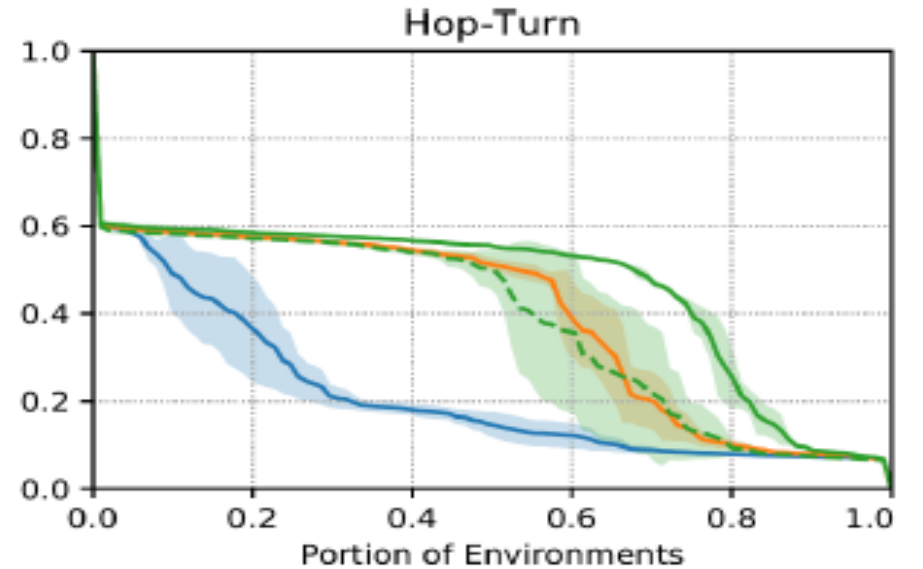
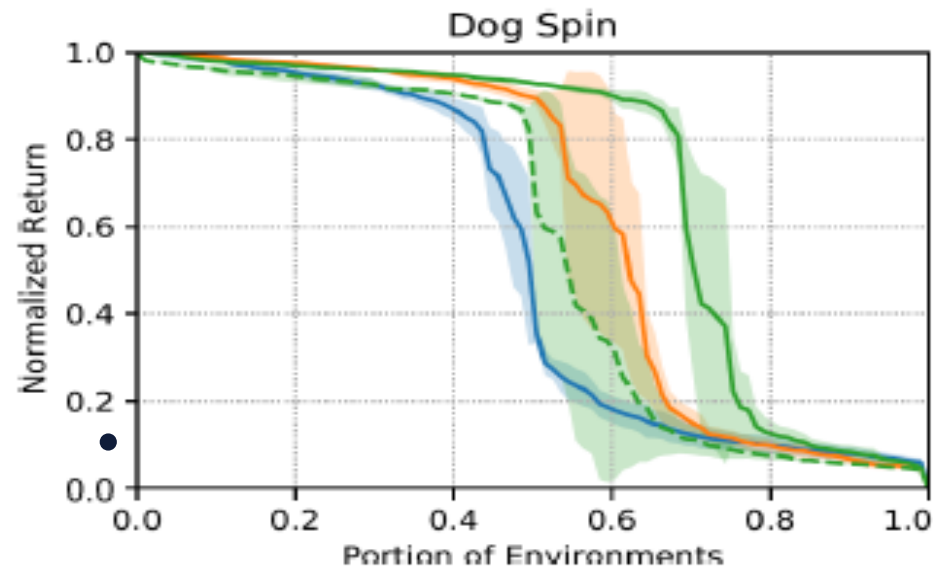
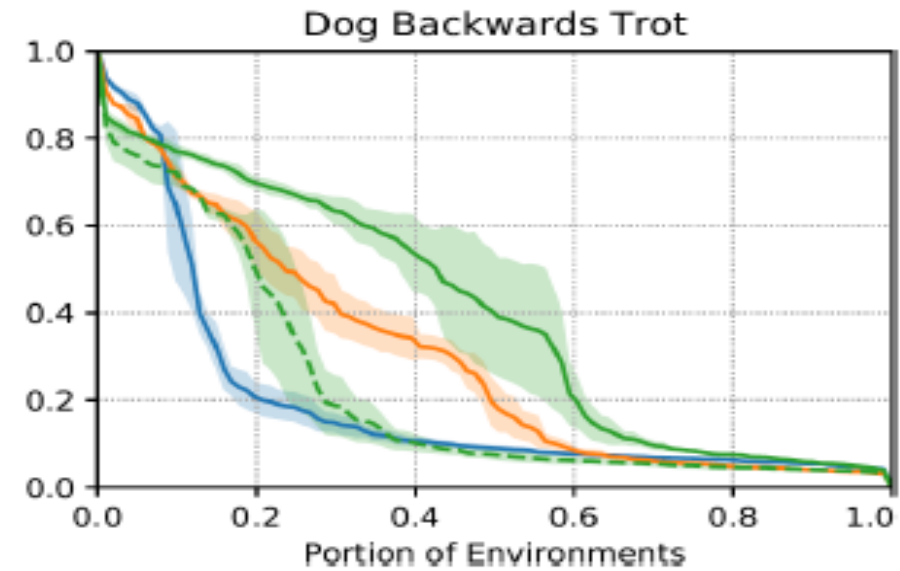
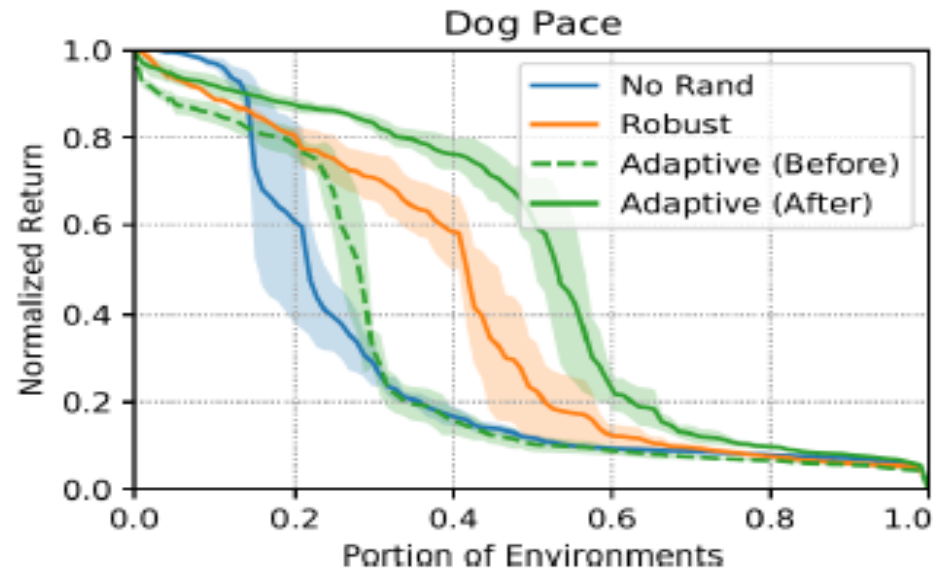
Experimental Setup

- Simulation performed on PyBullet
- MoCap clips from dog and artist generated renditions.
- Skills learned
 - Dog Pace
 - Dog Backwards Trot
 - Side-Steps
 - Turn
 - Hop-Turn

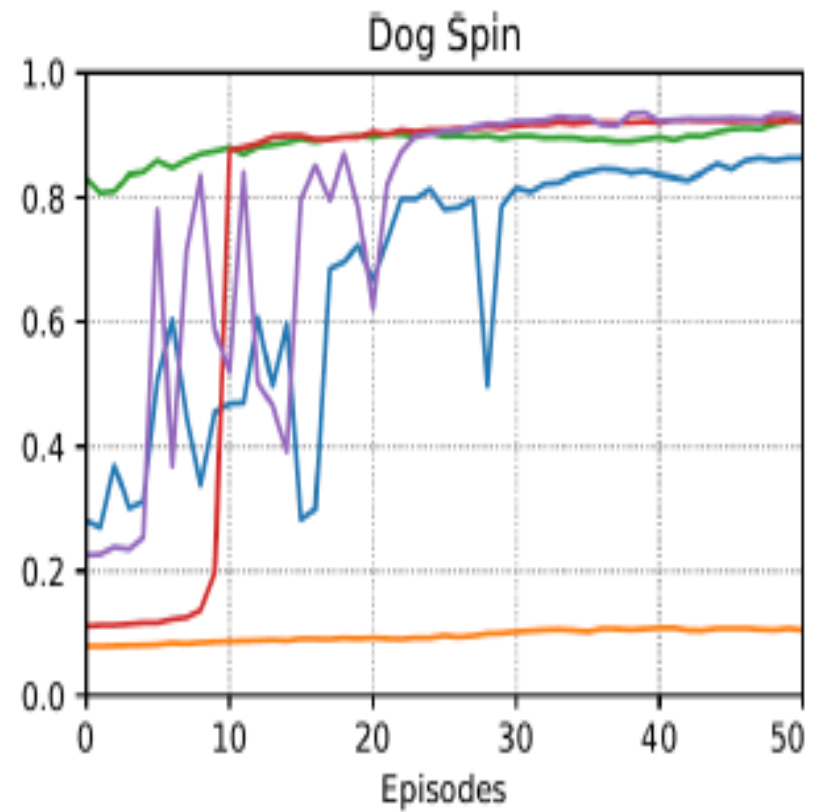
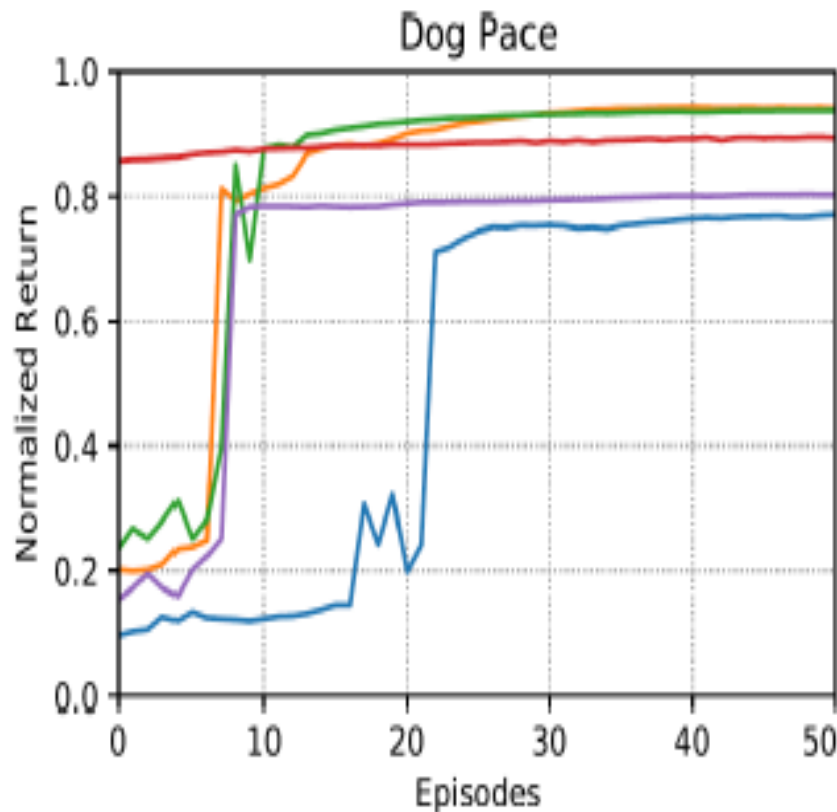
Policy Network



Results and Discussion

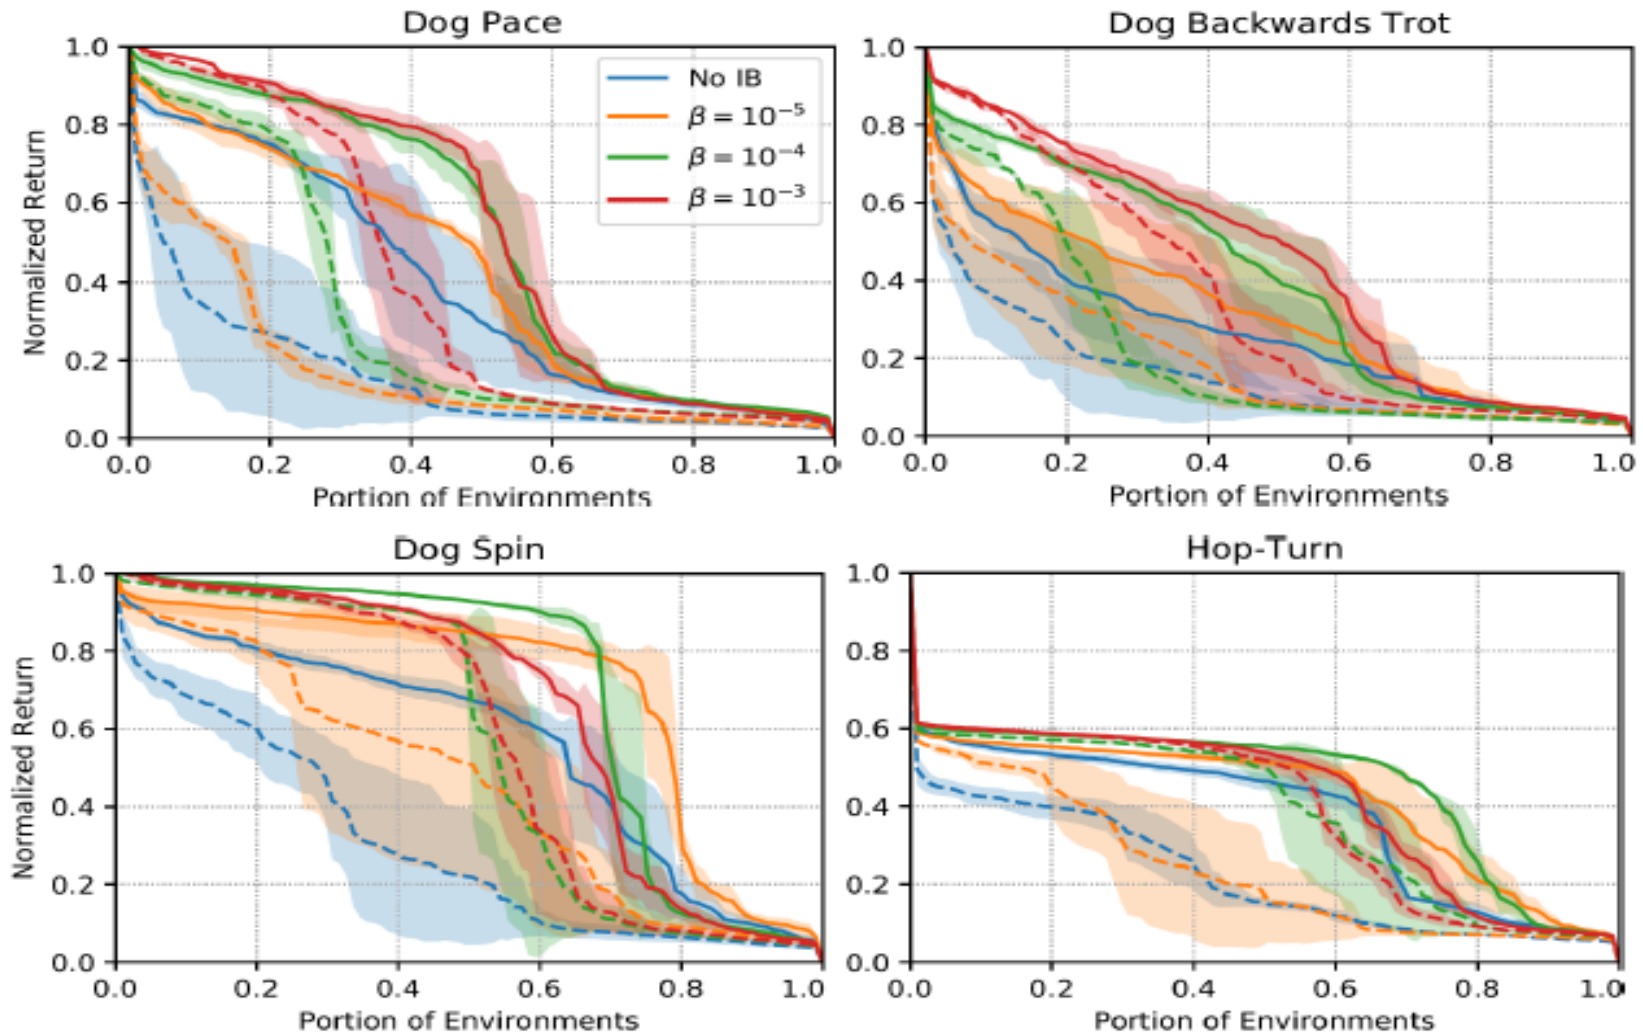


Results and Discussion

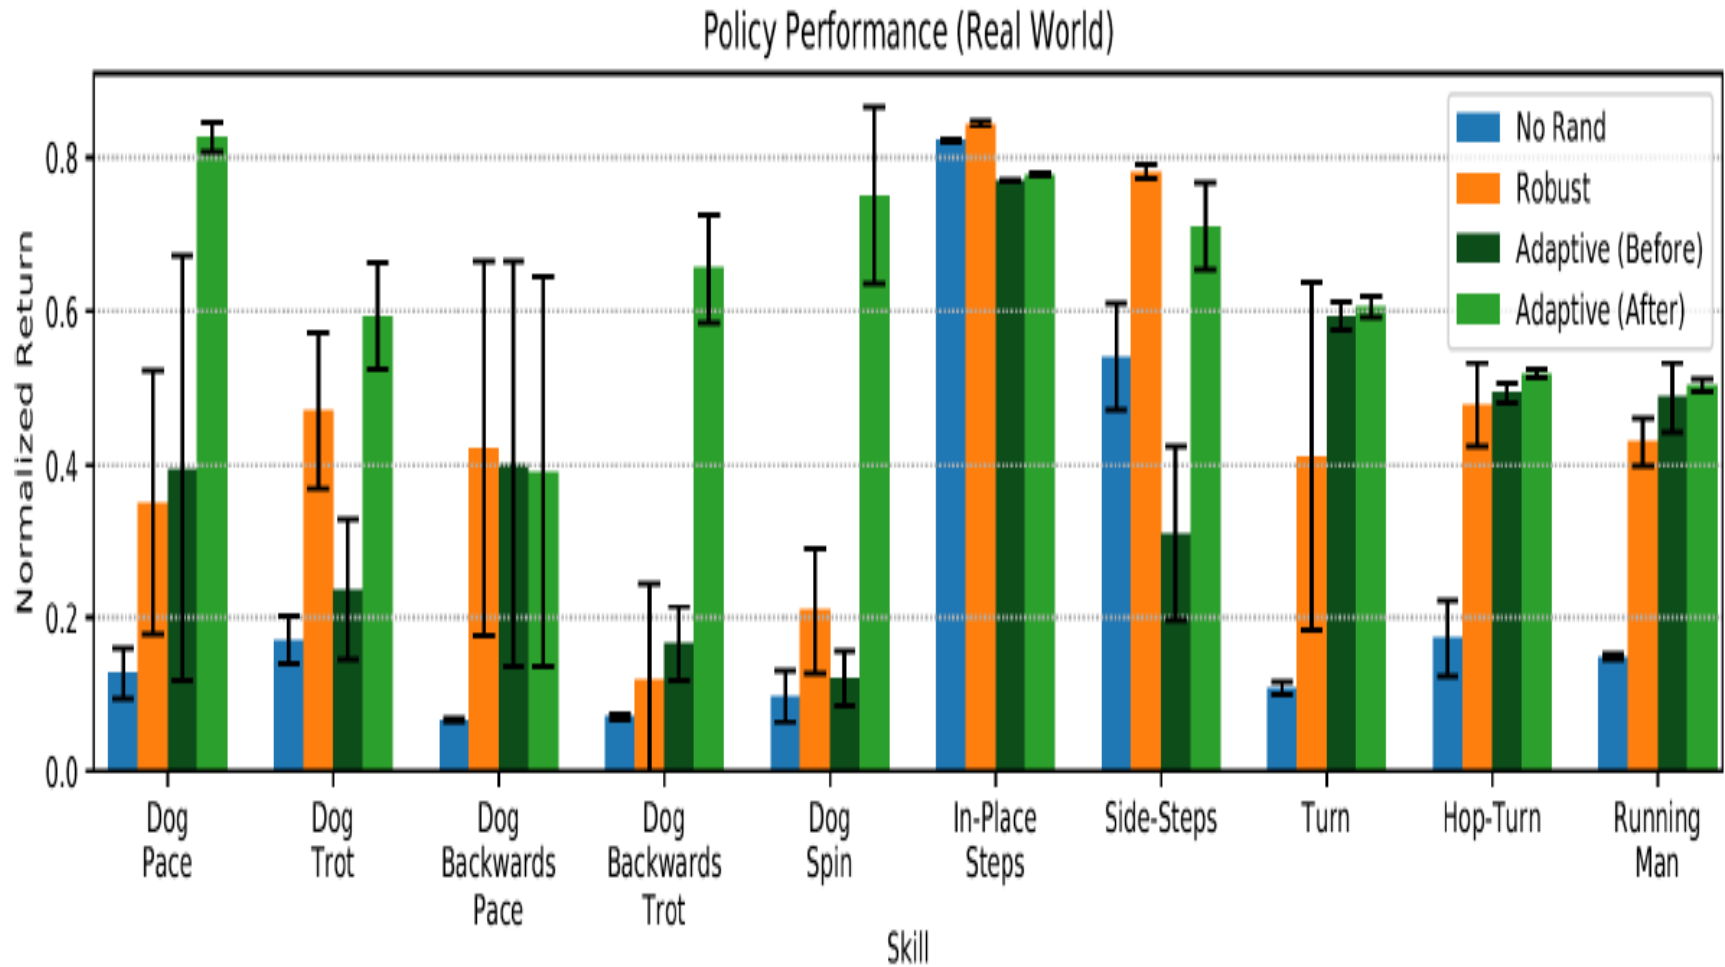


- Faster adaptation to new simulated environments

Results and Discussion



Results and Discussion



Results and Discussion

- Needless to say, the adaptive policies perform way better
- Experiments were conducted to measure normalised return of policies trained without randomisation, with randomisation, and with adaptation.
- Adaptive policies perform better for all kinds of motions.

Conclusion and Future Work

- The behaviours learned by our policies are currently not as stable as the best manually-designed controllers
- More robust controllers
- Learning from other sources of data, like videos

Thanks!!!