

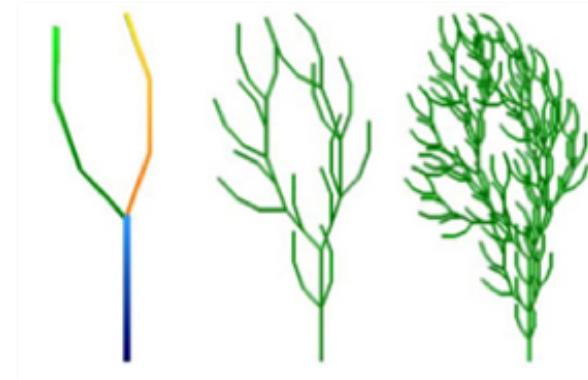
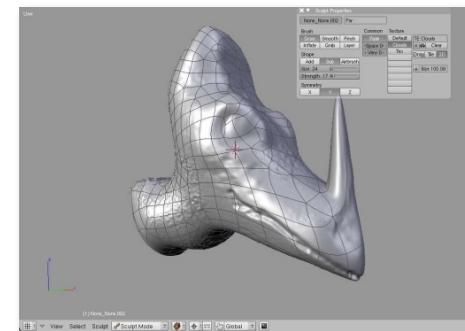
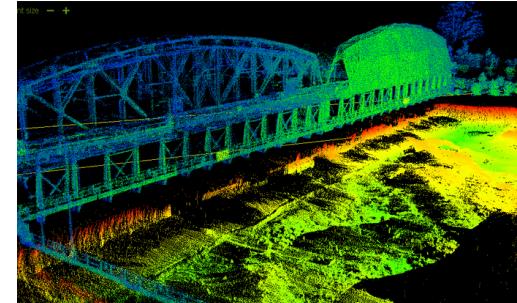
L4: Mesh and Point Cloud

Hao Su

We will
deviate from theory,
and
things will get dirty...


Shape Representation: Origin- and Application-Dependent

- Acquired real-world objects:
- Modeling “by hand”:
- Procedural modeling
- ...

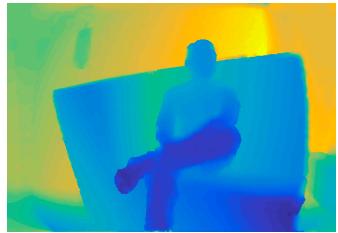


Other than parametric representations, we also study these in this course:

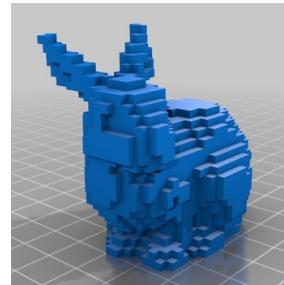
Rasterized form (regular grids)



Multi-view

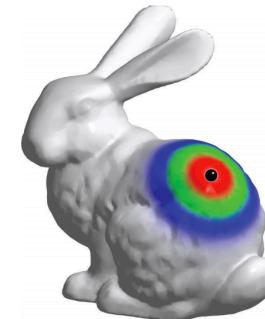


Depth Map

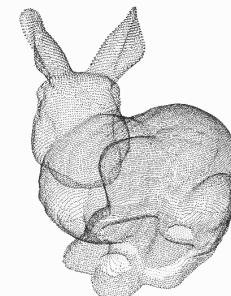


Volumetric

Geometric form (irregular)



Mesh

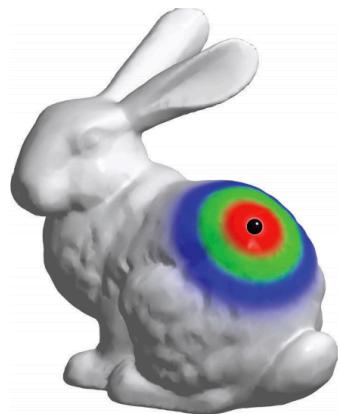


Point Cloud

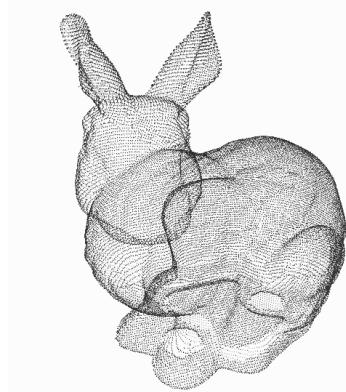
$$F(x) = 0$$

Implicit Shape

Agenda



Mesh



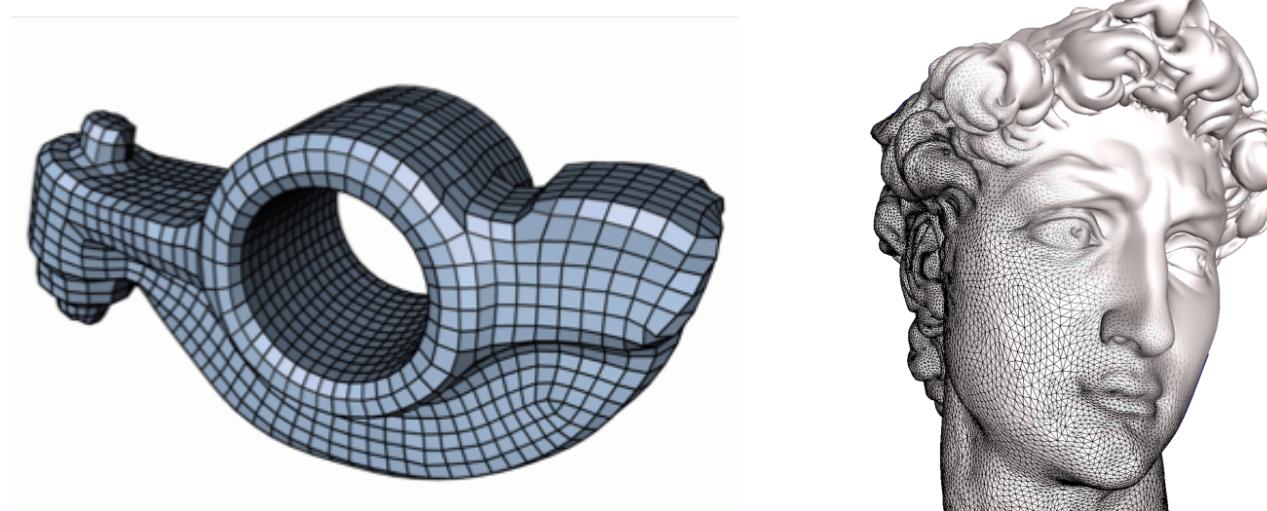
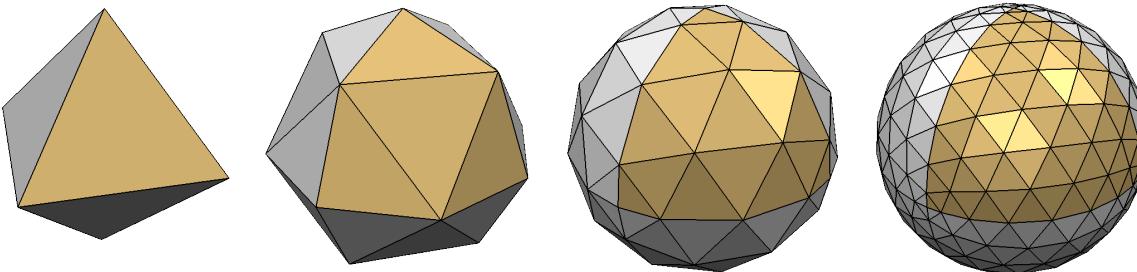
Point Cloud

Polygonal Meshes

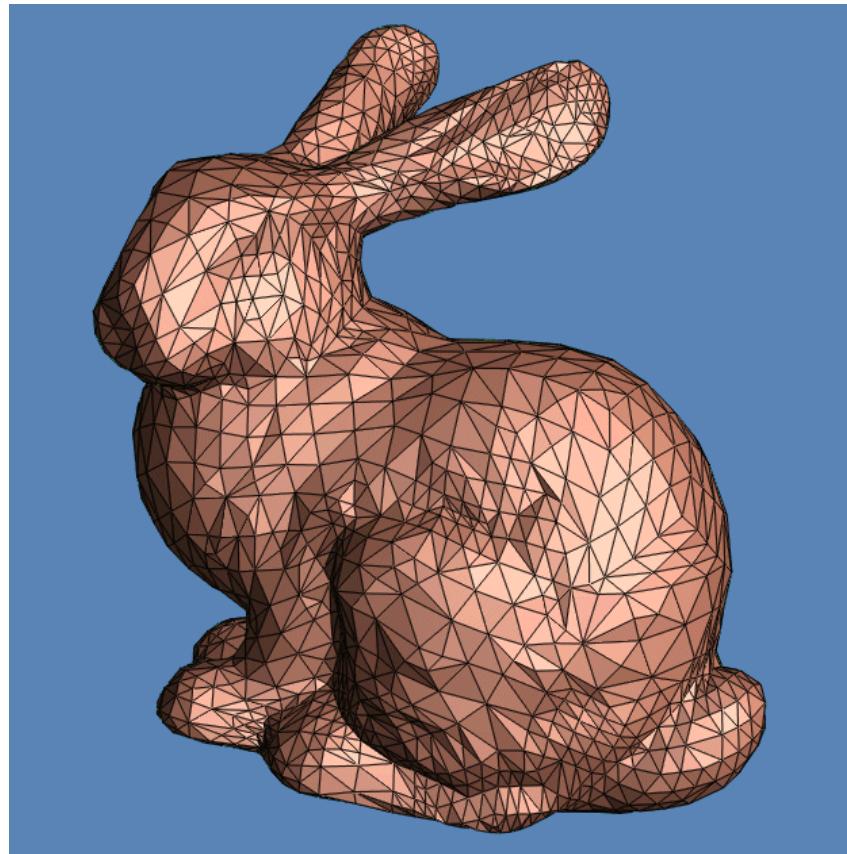
- Representation
- Storage
- Curvature Computation

Polygonal Meshes

- Boundary representations of objects



Triangle Mesh



<http://graphics.stanford.edu/data/3Dscanrep/stanford-bunny-cebal-ssh.jpg>
<http://www.stat.washington.edu/wxs/images/BUNMID.gif>

Triangle Mesh

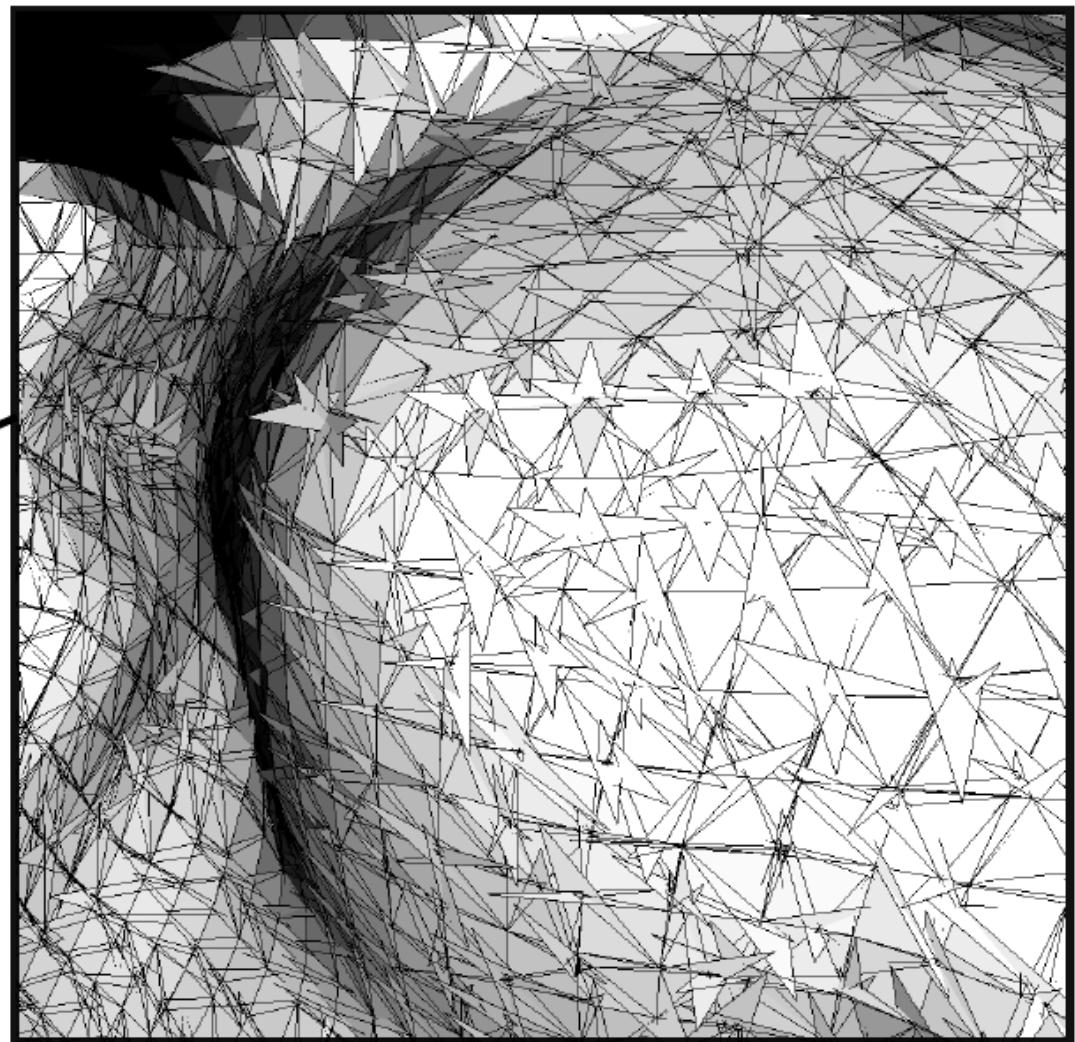
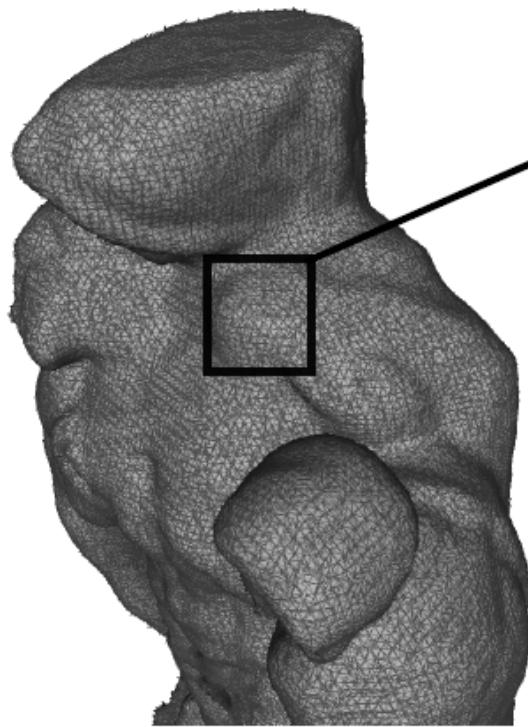
$$V = (v_1, v_2, \dots, v_n) \subset \mathbb{R}^n$$

$$E = (e_1, e_2, \dots, e_k) \subseteq V \times V$$


$$F = (f_1, f_2, \dots, f_m) \subseteq V \times V \times V$$

Plus manifold conditions

Bad Surfaces




Nonmanifold Edge




Manifold Mesh

1. Each **edge** is incident to one or two faces
2. **Faces** incident to a vertex form a closed or open fan

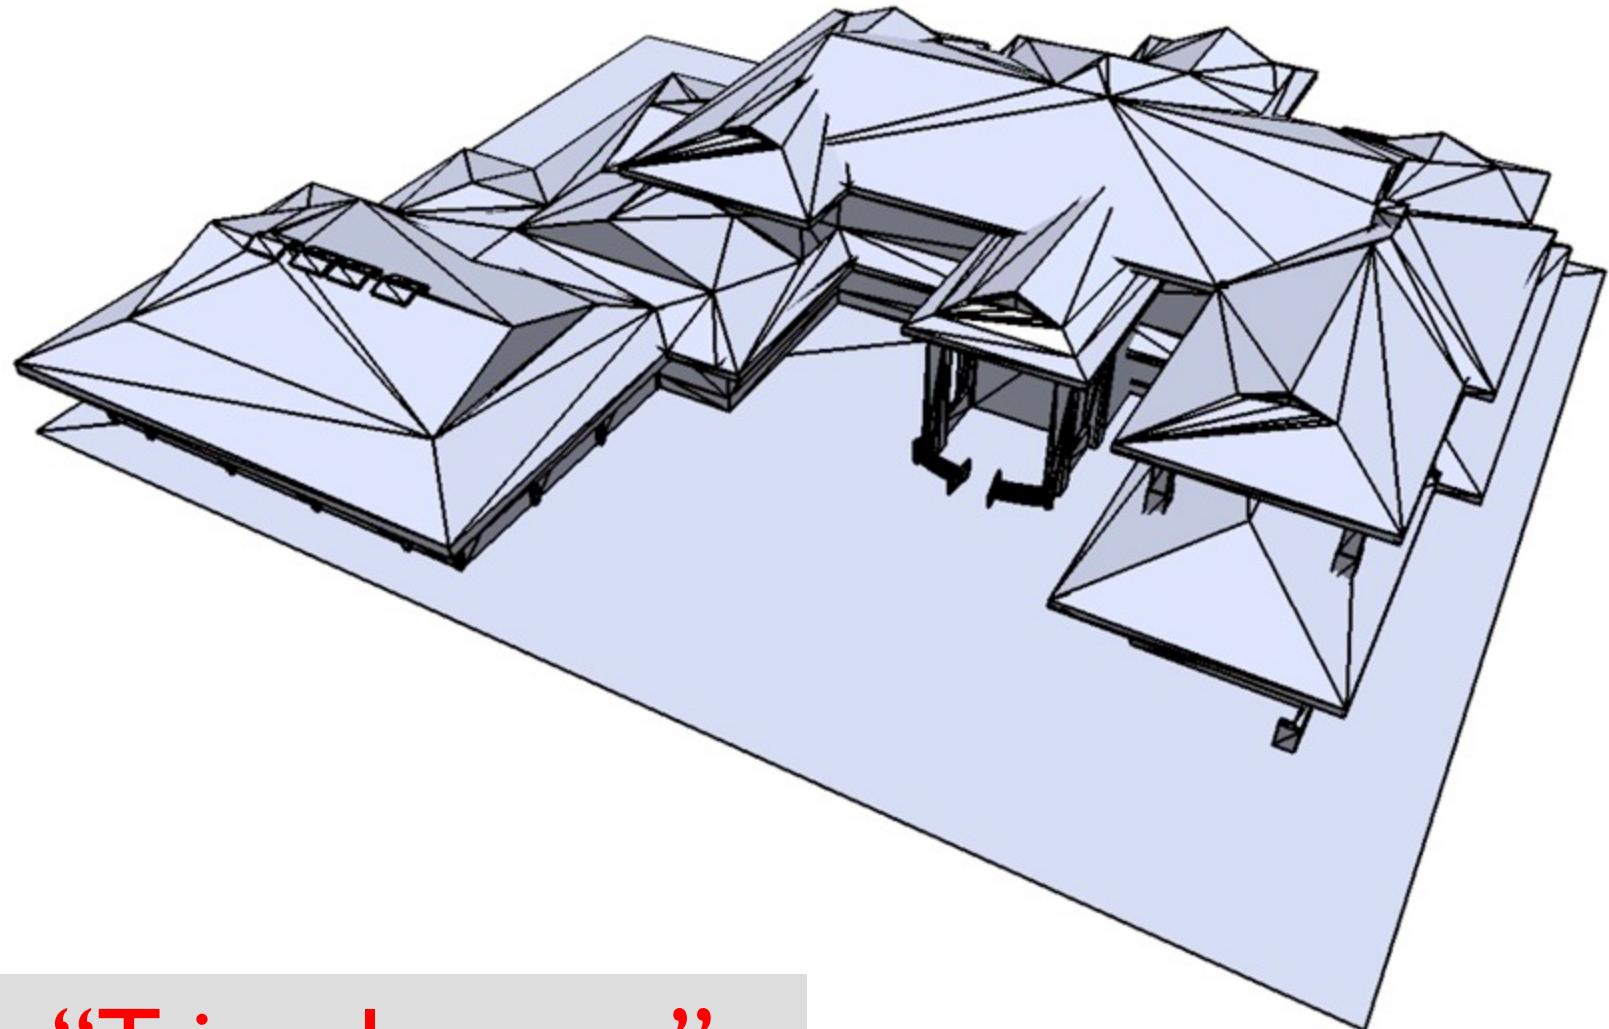


Manifold Mesh

1. Each **edge** is incident to one or two faces
2. **Faces** incident to a vertex form a closed or open fan




Assume meshes are manifold
(for now)



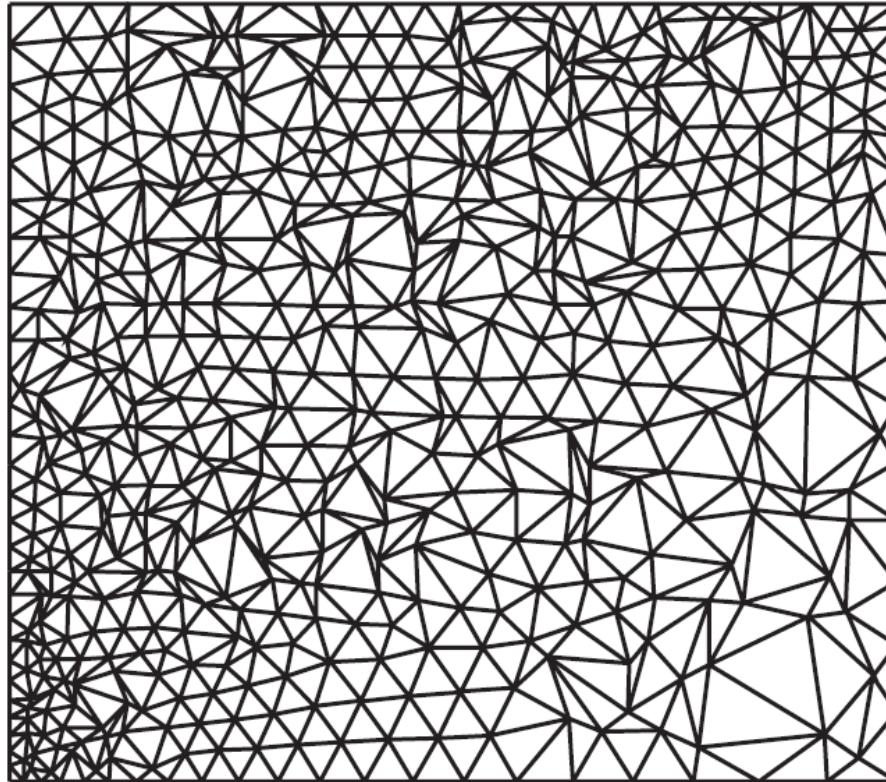
“Triangle soup”

Bad Meshes



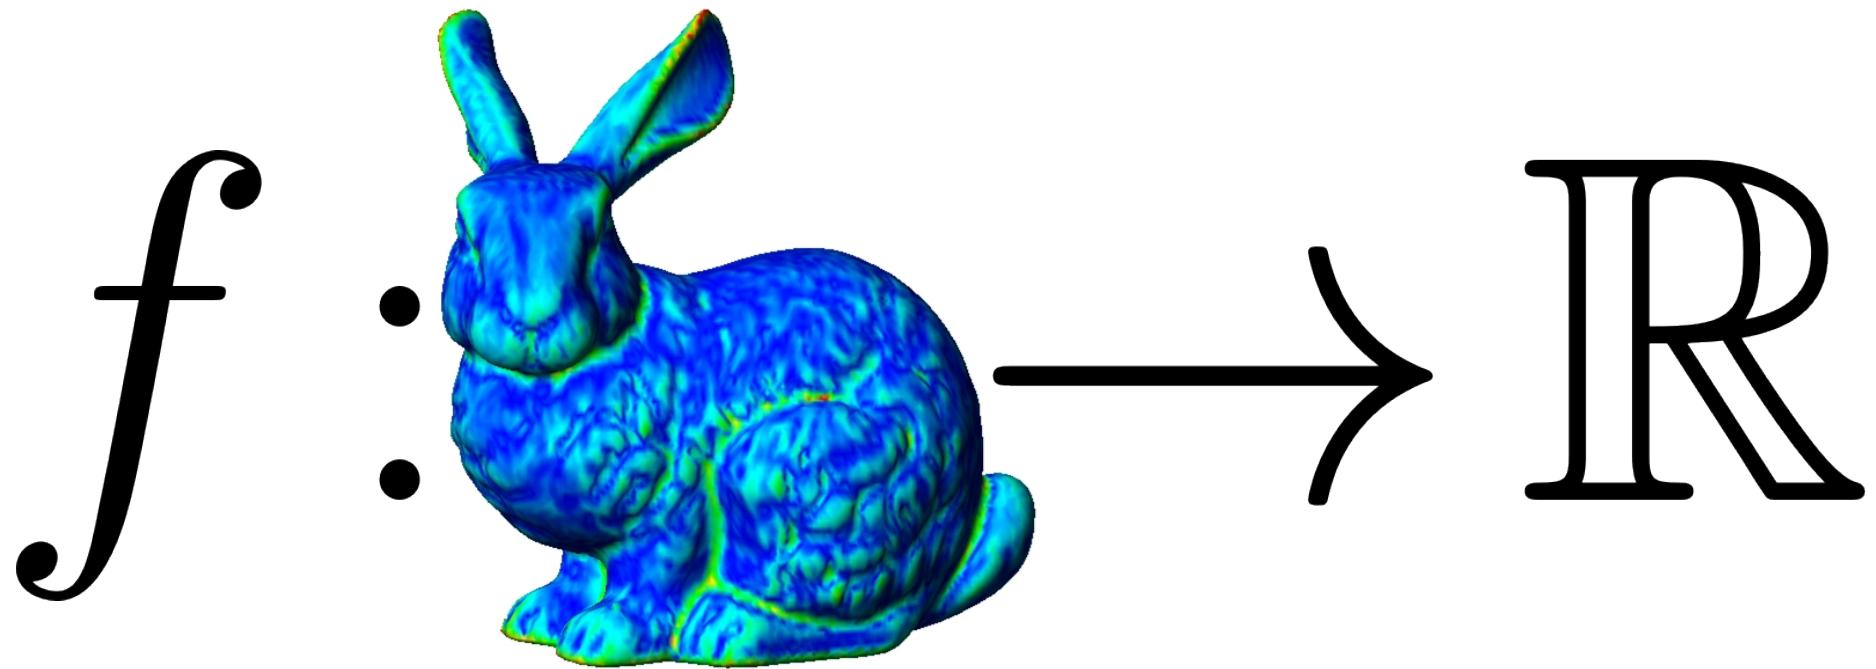
Nonuniform
areas and angles

Why is Meshing an Issue?



How do you interpret
one value per vertex?

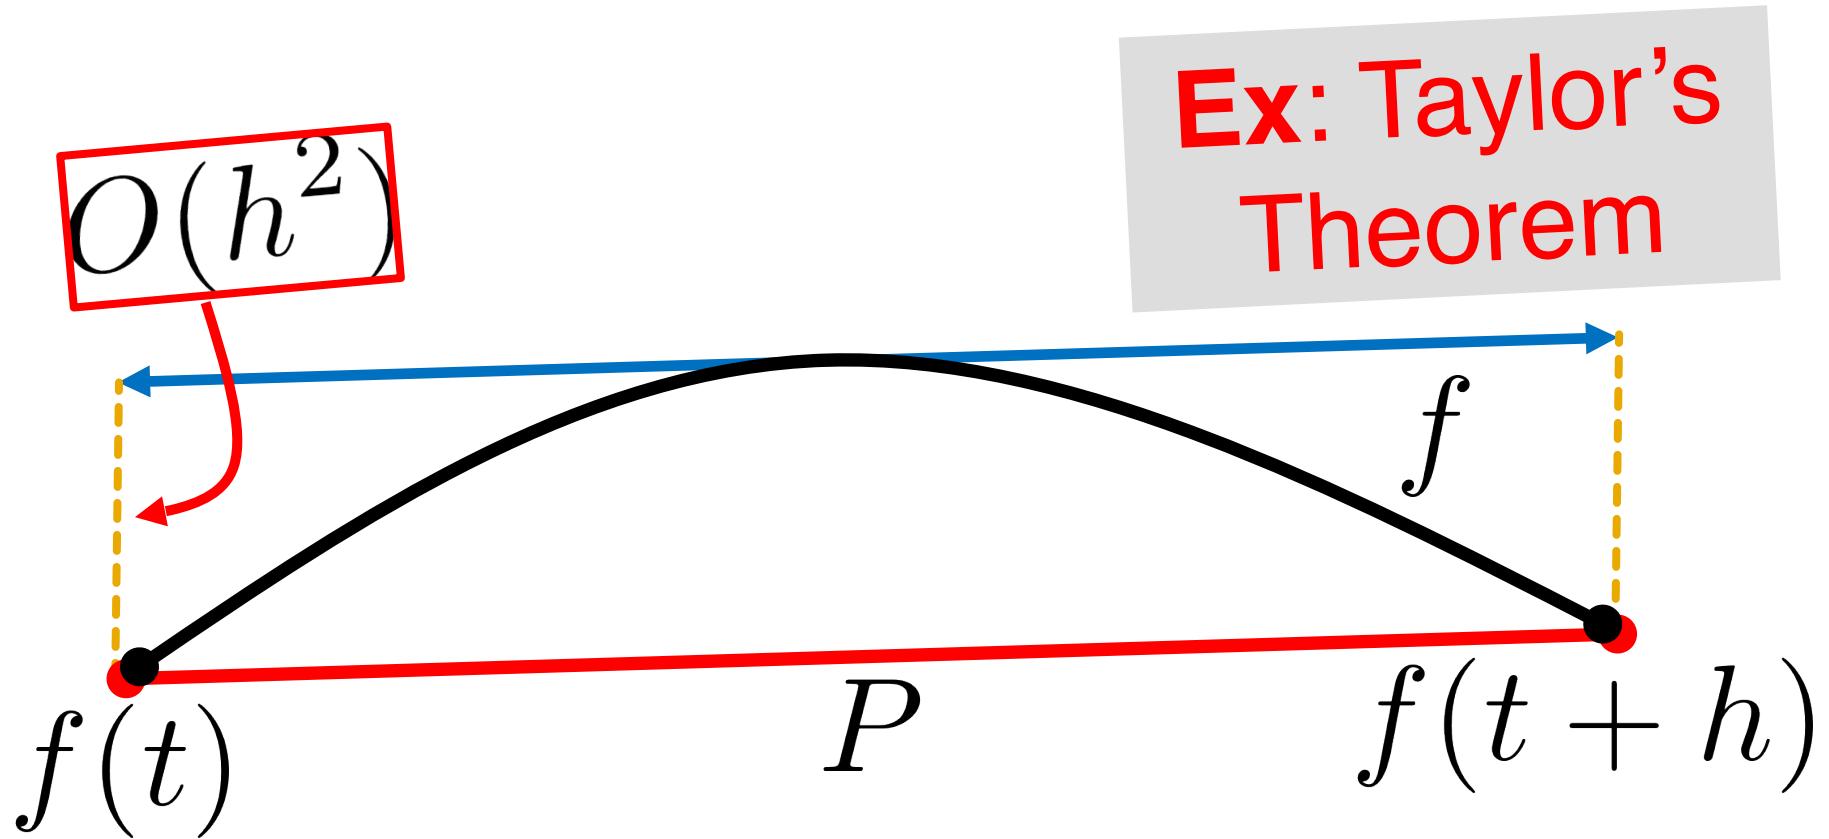
Assume Storing Scalar Functions on Surface



http://www.ieeta.pt/polymeco/Screenshots/PolyMeCo_OneView.jpg

Map points to real numbers


Approximation Properties



f : functions defined at vertices
(e.g., Gaussian curvature)

Techniques to Improve Mesh Quality

- Cleaning
- Repairing
- Remeshing
- ...



Polygonal Meshes

- Representation
- Storage
- Curvature Computation

Data Structures for Surfaces



- What should be stored?
 - Geometry: 3D coordinates
 - Topology
 - Attributes
 - Normal, color, texture coordinates
 - Per vertex, face, edge

Simple Data Structures: Triangle List

- STL format (used in CAD)
- Storage
 - Face: 3 positions
- No connectivity information

Triangles			
0	x0	y0	z0
1	x1	x1	z1
2	x2	y2	z2
3	x3	y3	z3
4	x4	y4	z4
5	x5	y5	z5
6	x6	y6	z6
...


Simple Data Structures: Indexed Face Set

- Used in formats
 - OBJ, OFF, WRL
- Storage
 - Vertex: position
 - Face: vertex indices
 - Convention is to save vertices in counter-clockwise order for normal computation

Vertices			
v0	x0	y0	z0
v1	x1	x1	z1
v2	x2	y2	z2
v3	x3	y3	z3
v4	x4	y4	z4
v5	x5	y5	z5
v6	x6	y6	z6
...

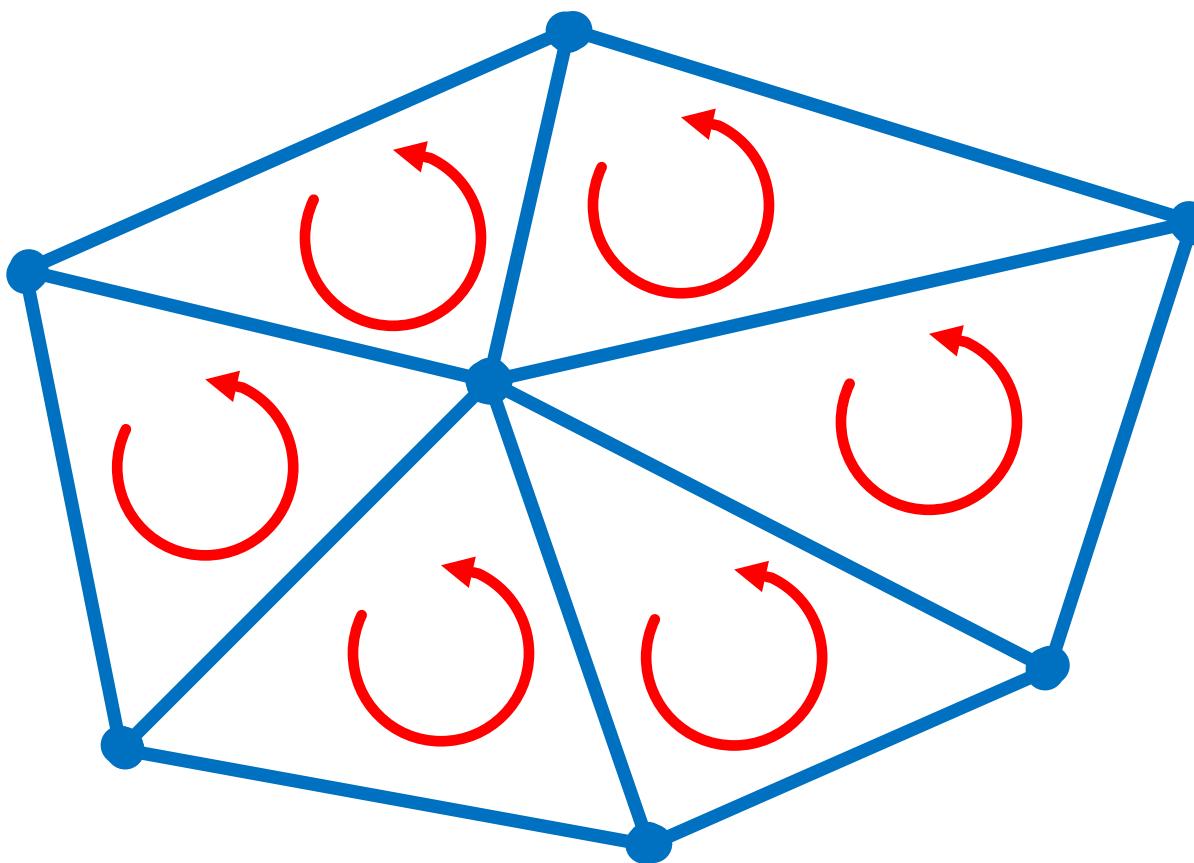
Triangles			
t0	v0	v1	v2
t1	v0	v1	v3
t2	v2	v4	v3
t3	v5	v2	v6
...

Right-Hand Rule




http://viz.aset.psu.edu/gho/sem_notes/3d_fundamentals/html/3d_coordinates.html
http://mathinsight.org/stokes_theorem_orientation

Normal Computation



Orientability



orientable



non-orientable


Summary of Polygonal Meshes

- Polygonal meshes are piece-wise linear approximation of smooth surfaces
- Good triangulation is important (manifold, equi-length)
- Vertices, edges, and faces are basic elements
- While real-data 3D are often point clouds, meshes are quite often used to visualize 3D and generate ground truth for machine learning algorithms

Polygonal Meshes


- Representation
- Storage
- Curvature Computation

Challenge on Meshes



Curvature is a
second derivative,
but triangles are flat.

Rusinkiewicz's Method



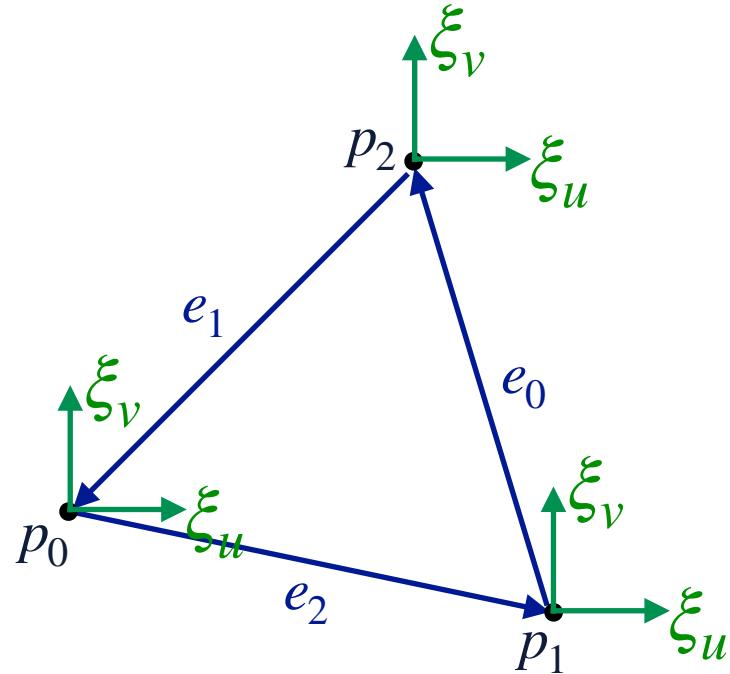
Assume a local $f: U \rightarrow \mathbb{R}^3$ at a small triangle

Assume that \mathbf{T}_{p_i} 's are roughly aligned

Assume that $Df \begin{bmatrix} u \\ v \end{bmatrix} = u \vec{\xi}_u + v \vec{\xi}_v$, i.e., $Df = [\vec{\xi}_u, \vec{\xi}_v]$

(U is also aligned with \mathbf{T}_{p_i} 's)

Rusinkiewicz's Method



Recall shape operator: $DN = Df \cdot S$.

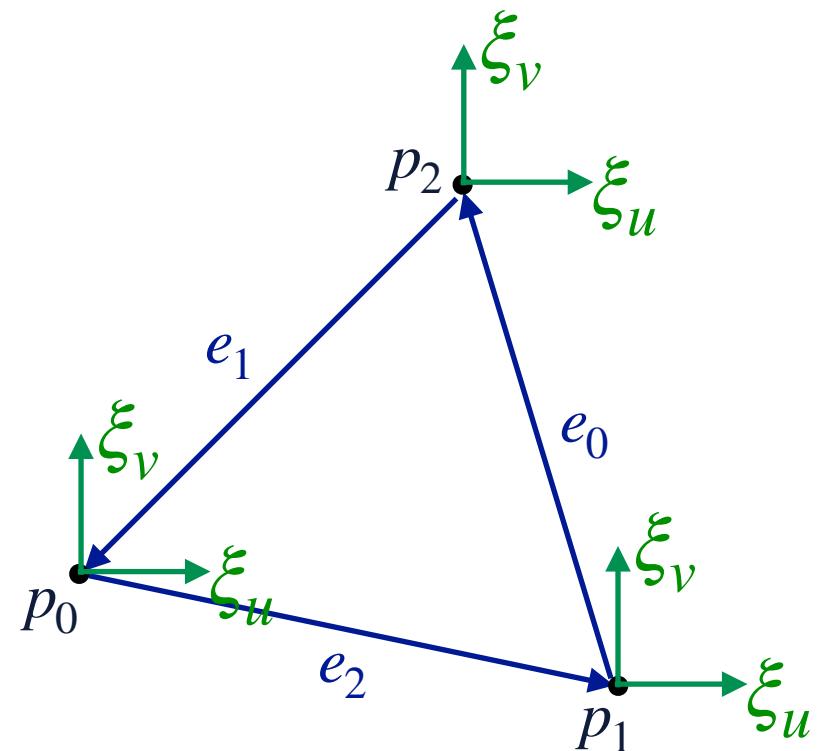
$$\therefore Df = \left[\vec{\xi}_u, \vec{\xi}_v \right], \therefore S = Df^T DN$$

If we have S , we can compute principal curvatures!
How to estimate S ?

$$Df^T \left(DN \begin{bmatrix} u \\ v \end{bmatrix} \right) \approx Df^T \Delta \vec{n} \implies S \begin{bmatrix} u \\ v \end{bmatrix} \approx Df^T \Delta \vec{n}$$

$$\because Df \begin{bmatrix} u \\ v \end{bmatrix} = Y \in \mathbf{T}(\mathbb{R}^3) \text{ and } Df = [\vec{\xi}_u, \vec{\xi}_v] \quad \therefore \begin{bmatrix} u \\ v \end{bmatrix} = Df^T Y$$

$$\therefore S[Df]^T \begin{bmatrix} u \\ v \end{bmatrix} \approx [Df]^T \Delta \vec{n}^T$$




$$Df^T \left(DN \begin{bmatrix} u \\ v \end{bmatrix} \right) \approx Df^T \Delta \vec{n} \implies S \begin{bmatrix} u \\ v \end{bmatrix} \approx Df^T \Delta \vec{n}$$

$$\because Df \begin{bmatrix} u \\ v \end{bmatrix} = Y \in \mathbf{T}(\mathbb{R}^3) \text{ and } Df = [\vec{\xi}_u, \vec{\xi}_v] \quad \therefore \begin{bmatrix} u \\ v \end{bmatrix} = Df^T Y$$

$$\therefore S[Df]^T \begin{bmatrix} u \\ v \end{bmatrix} \approx [Df]^T \Delta \vec{n}^T$$

$$\begin{cases} S[Df]^T e_0 = Df^T (\vec{n}_2 - \vec{n}_1), \\ S[Df]^T e_1 = Df^T (\vec{n}_0 - \vec{n}_2), \\ S[Df]^T e_2 = Df^T (\vec{n}_1 - \vec{n}_0), \end{cases}$$

So we can solve $S \in \mathbb{R}^{2 \times 2}$ by least square (6 equations and 4 unknowns)



Summary of Mesh Curvature Estimation

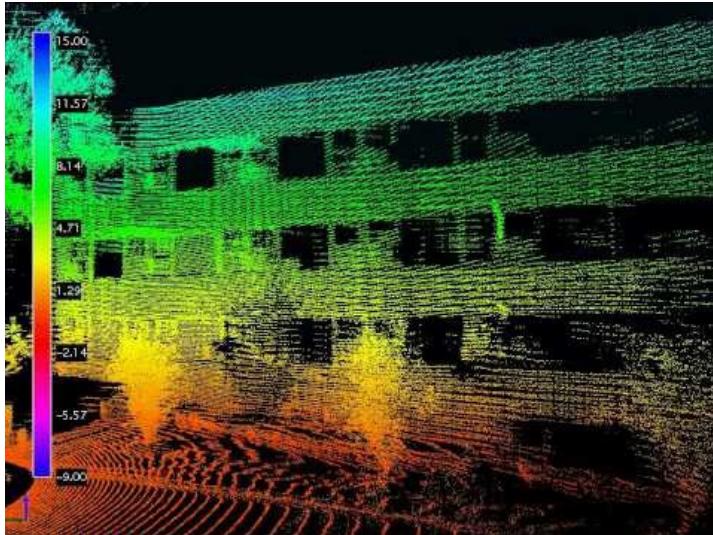
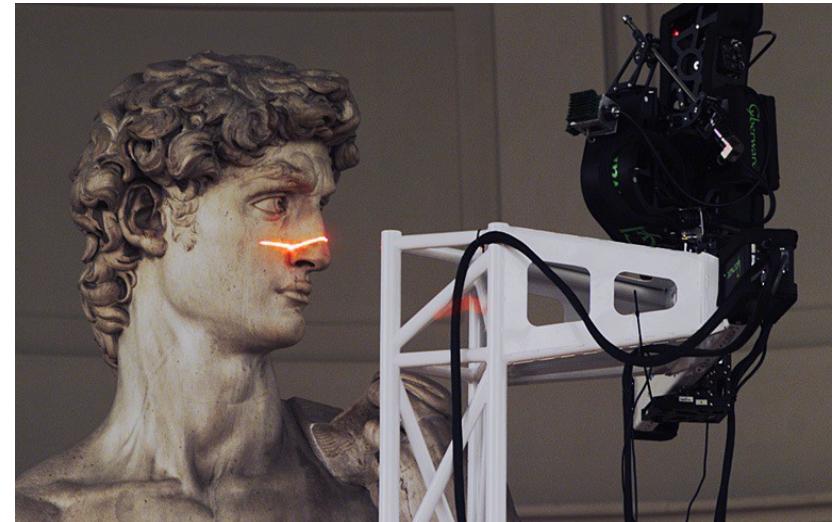
- Rusinkiewicz's method is an effective approach for face curvature estimation
 - Szymon Rusinkiewicz, "Estimating Curvatures and Their Derivatives on Triangle Meshes", 3DPVT, 2004
- Good robustness to moderate amount of noise and free of degenerate configurations
- Can be used to compute curvatures for point cloud as well

Point Cloud

- Representation
- Sampling Points on Surfaces
- Normal Computation

Acquiring Point Clouds

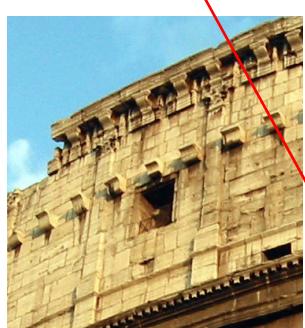
- From the real world
 - 3D scanning
 - Data is “striped”
 - Need multiple views to compensate occlusion



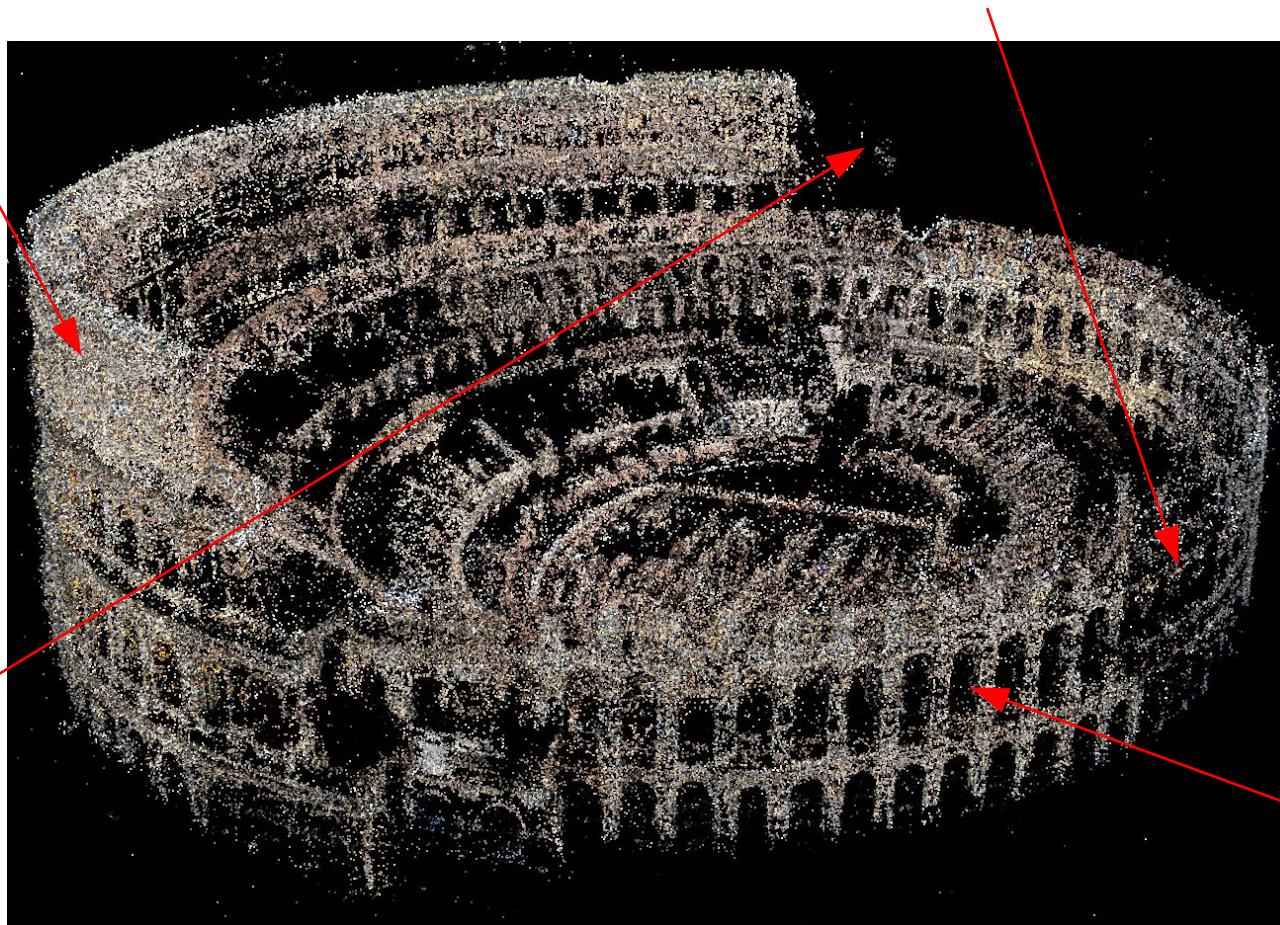
- Many techniques
 - Laser (LIDAR, e.g., StreetView)
 - Infrared (e.g., Kinect)
 - Stereo (e.g., Bundler)
- Many challenges: resolution, occlusion, noise, registration

Acquisition Challenges

Noise → Poor detail reproduction



Low resolution further obscures detail

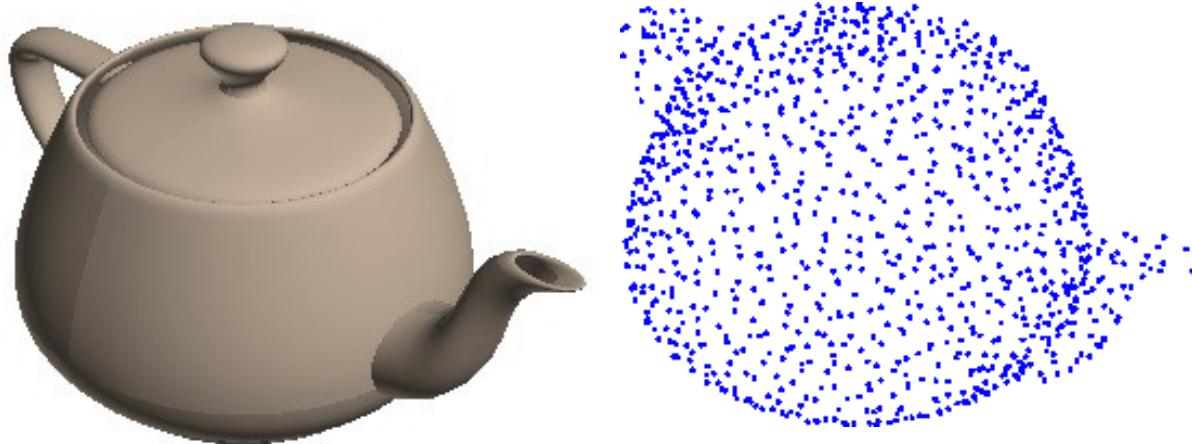


Some data was
not properly
registered with
the rest

Occlusion →
Interiors not
captured

Acquiring Point Clouds

- From existing virtual shapes



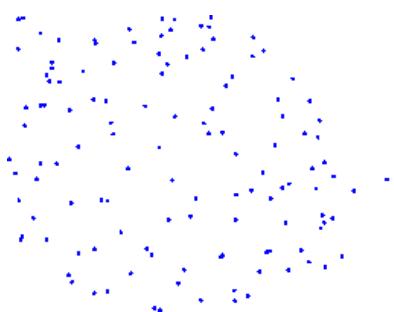
- Why would we want to do this?

Light-weight Shape Representation

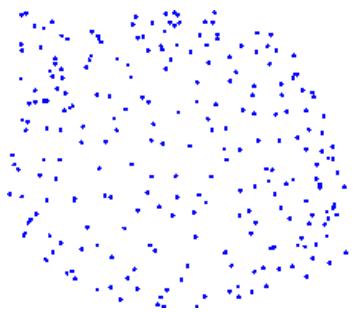
Point cloud:

- Simple to understand
- Compact to store
- Generally easy to build algorithms

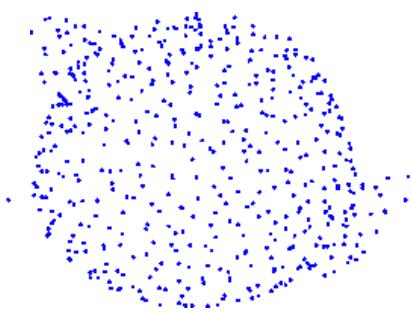
Yet already carries rich information!



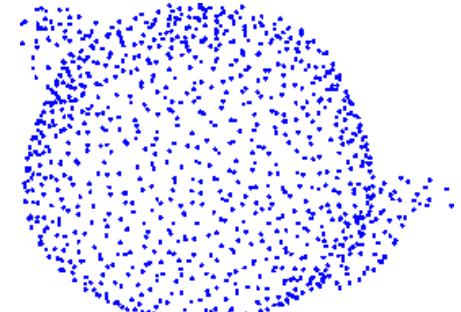
$N = 125$



$N = 250$



$N = 500$



$N = 1000$

Point Cloud

- Representation
- Sampling Points on Surfaces
- Normal Computation

Application-based Sampling

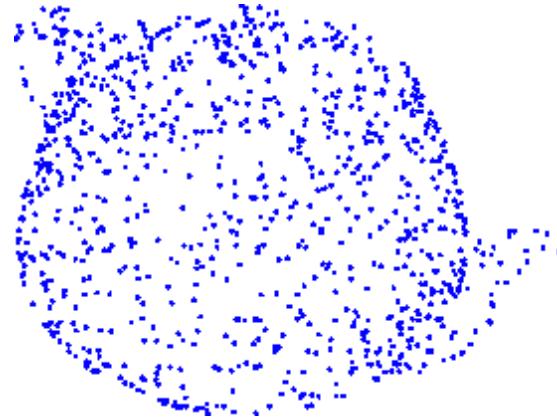
- For storage or analysis purposes (e.g., shape retrieval, signature extraction),
 - the objective is often to preserve surface information as much as possible
- For learning data generation purposes (e.g., sim2real),
 - the objective is often to minimize virtual-real domain gap

Application-based Sampling

- For storage or analysis purposes (e.g., shape retrieval, signature extraction),
 - the objective is often to preserve surface information as much as possible
- For learning data generation purposes (e.g., sim2real),
 - the objective is often to minimize virtual-real domain gap

Naive Strategy: Uniform Sampling


- Independent identically distributed (i.i.d.) samples by surface area:



- Usually the easiest to implement (as in your HW0)
- Issue: Irregularly spaced sampling

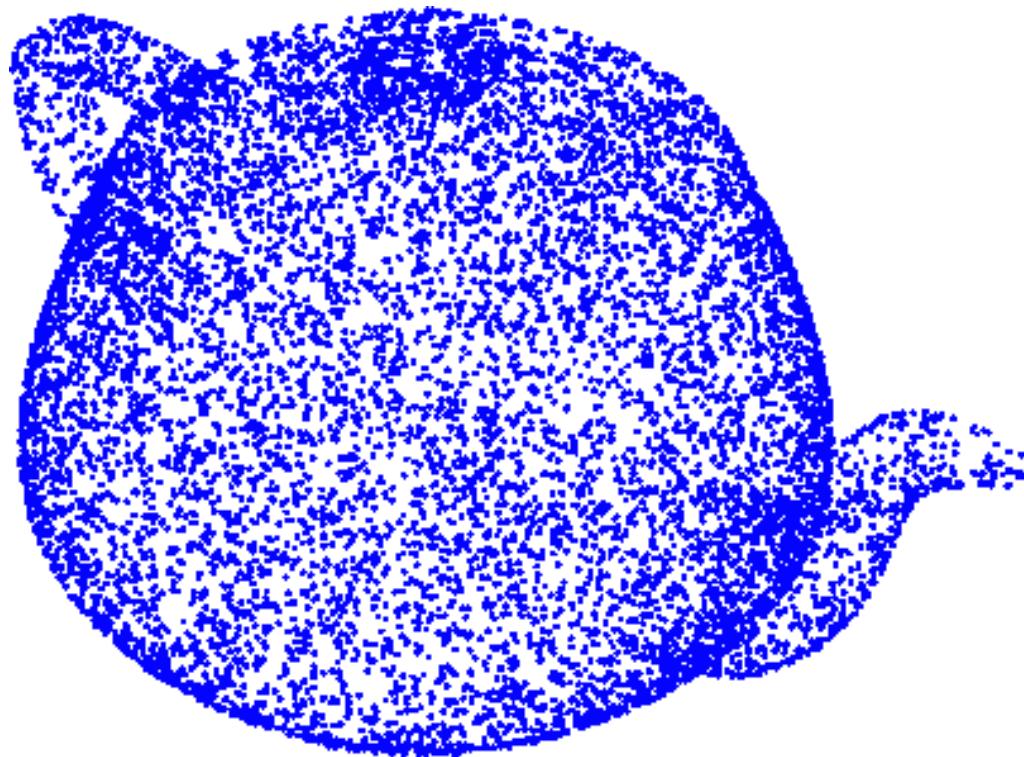
Farthest Point Sampling

- Goal: Sampled points are far away from each other
- NP-hard problem
- What is a greedy approximation method?



Iterative Furthest Point Sampling

- Step 1: Over sample the shape by any fast method
(e.g., uniformly sample $N=10,000$ i.i.d. samples)



Iterative Furthest Point Sampling

- Step 2: Iteratively select K points

U is the initial big set of points

$S = \{ \}$

add a random point from U to S

for $i=1$ to K

 find a point $u \in U$ with the largest distance to S

 add u to S

Issues Relevant to Speed

- Theoretically, naive implementation gives $\mathcal{O}(KN)$, but how to improve from $\mathcal{O}(KN)$ is an open question
- **Implementation can cause large speed difference**
 - As this is a serial algorithm in K , engineers optimize the efficiency in N (computing point-set distance)
 - CPU: Suggest using vectorization (e.g., numpy, `scipy.spatial.distance.cdist`)
 - GPU: By using shared memory, the complexity can be reduced to $\mathcal{O}(K(N/M + \log M))$, where M is the number of threads ($M=512$ in practice for modern GPU).

Implementation Tricks

- References:
 - https://github.com/maxjaritz/mvpnet/blob/master/mvpnet/ops/cuda/fps_kernel.cu
 - https://github.com/erikwijmans/Pointnet2_PyTorch/blob/master/pointnet2_ops_lib/pointnet2_ops/ext-src/src/sampling_gpu.cu
- By courtesy of Jiayuan Gu, we share a GPU version code with you (through Piazza)

Read by yourself!

An Implementation in Numpy



```
def fps_downsample(points, number_of_points_to_sample):
    selected_points = np.zeros((number_of_points_to_sample, 3))
    dist = np.ones(points.shape[0]) * np.inf # distance to the selected set
    for i in range(number_of_points_to_sample):
        # pick the point with max dist
        idx = np.argmax(dist)
        selected_points[i] = points[idx]
        dist_ = ((points - selected_points[i]) ** 2).sum(-1)
        dist = np.minimum(dist, dist_)

    return selected_points
```

Read by yourself!

Voxel Downsampling

- Uses a *regular* voxel grid to downsample, taking one point per grid
- Allows higher parallelization level
- Generates regularly spaced sampling (with noticeable artifacts)



Issues Relevant to Speed

- Need to map each point to a bin. Often implemented as adding elements into a hash table
- $\mathcal{O}(N)$ (assuming that the inserting into hash table takes $O(1)$)
- On GPUs, parallelization reduces complexity of
 - Mapping each point to an integer value
 - Assign each value to an index so that the same value shares the same index
 - Aggregate indexes and form the output (called scattering in CUDA)

A Hashtable-based Implementation in Numpy

```
def voxel_downsample(points: np.ndarray, voxel_size: float):
    """Voxel downsample (first).

    Args:
        points: [N, 3]
        voxel_size: scalar

    Returns:
        np.ndarray: [M, 3]
    """
    points_downsampled = dict() # point in each voxel cell
    points_voxel_coords = (points / voxel_size).astype(int) # discretize to voxel
    coordinate
    for point_idx, voxel_coord in enumerate(points_voxel_coords):
        key = tuple(voxel_coord.tolist()) # voxel coordinate
        if key not in points_downsampled:
            # assign the point to a voxel cell
            points_downsampled[key] = points[point_idx]
    points_downsampled = np.array(list(points_downsampled.values()))
    return points_downsampled
```

A Hashtable-based Implementation in Torch

```
def voxel_downsample_torch(points: torch.Tensor, voxel_size: float):
    """Voxel downsample (average).

    Args:
        points: [N, 3]
        voxel_size: scalar

    Returns:
        torch.Tensor: [M, 3]
    """
    points = torch.as_tensor(points, dtype=torch.float32)
    points_voxel_coords = (points / voxel_size).long()  # discretize

    # Generate the assignment between points and voxel cells
    unique_voxel_coords, points_voxel_indices, count_voxel_coords = torch.unique(
        points_voxel_coords, return_inverse=True, return_counts=True, dim=0)

    M = unique_voxel_coords.size(0)  # the number of voxel cells
    points_downsampled = points.new_zeros([M, 3])
    points_downsampled.scatter_add_(
        dim=0,
        index=points_voxel_indices.unsqueeze(-1).expand(-1, 3),
        src=points)
    points_downsampled = points_downsampled / count_voxel_coords.unsqueeze(-1)
    return points_downsampled
```

Read by yourself!

Application-based Sampling

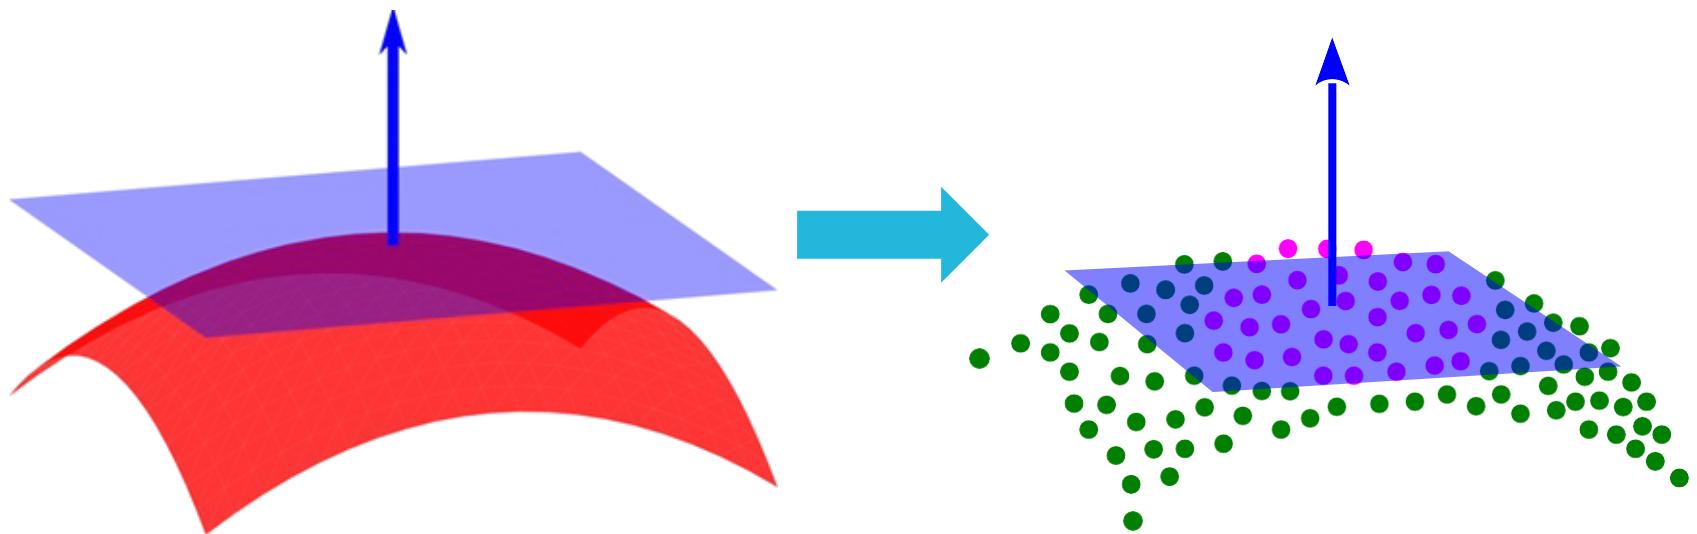
- For storage or analysis purposes (e.g., shape retrieval, signature extraction),
 - the objective is often to preserve surface information as much as possible
- For learning data generation purposes (e.g., sim2real),
 - the objective is often to minimize virtual-real domain gap
 - **a good research topic (e.g., GAN? Adversarial training? Differentiable sampling?)**

Point Cloud

- Representation
- Sampling Points on Surfaces
- Normal Computation

Estimating Normals

- Plane-fitting: find the plane that best fits the neighborhood of a point of interest



Least-square Formulation

- Assume the plane equation is:

$$w^T(x - c) = 0 \quad \text{with} \quad \|w\| = 1$$


- Plane-fitting solves the least square problem:

$$\begin{aligned} & \text{minimize} && \sum_i \|w^T(x_i - c)\|_2^2 \\ & \text{subject to} && \|w\|^2 = 1 \end{aligned}$$

where $\{x_i\}$ is the neighborhood of a point x
that you query the normal

$$\begin{array}{ll}\text{minimize} & \sum_i \|w^T(x_i - c)\|_2^2 \\ \text{subject to} & \|w\|^2 = 1\end{array}$$

- Doing Lagrangian multiplier and the solution is:
 - Let $M = \sum_i (x_i - \bar{x})(x_i - \bar{x})^T$ and $\bar{x} = \frac{1}{n} \sum_i x_i$,
 - w : the smallest eigenvector of M
 - $c = w^T \bar{x}$
- w also corresponds to the third principal component of $\{x_i\}$ (yet another usage of PCA)
 - Where are the first and second components?



Summary of Normal Computation

- The normal of a point cloud can be computed through PCA over a local neighborhood
- Remark:
 - The choice of neighborhood size is important
 - When outlier points exist, RANSAC can improve quality